



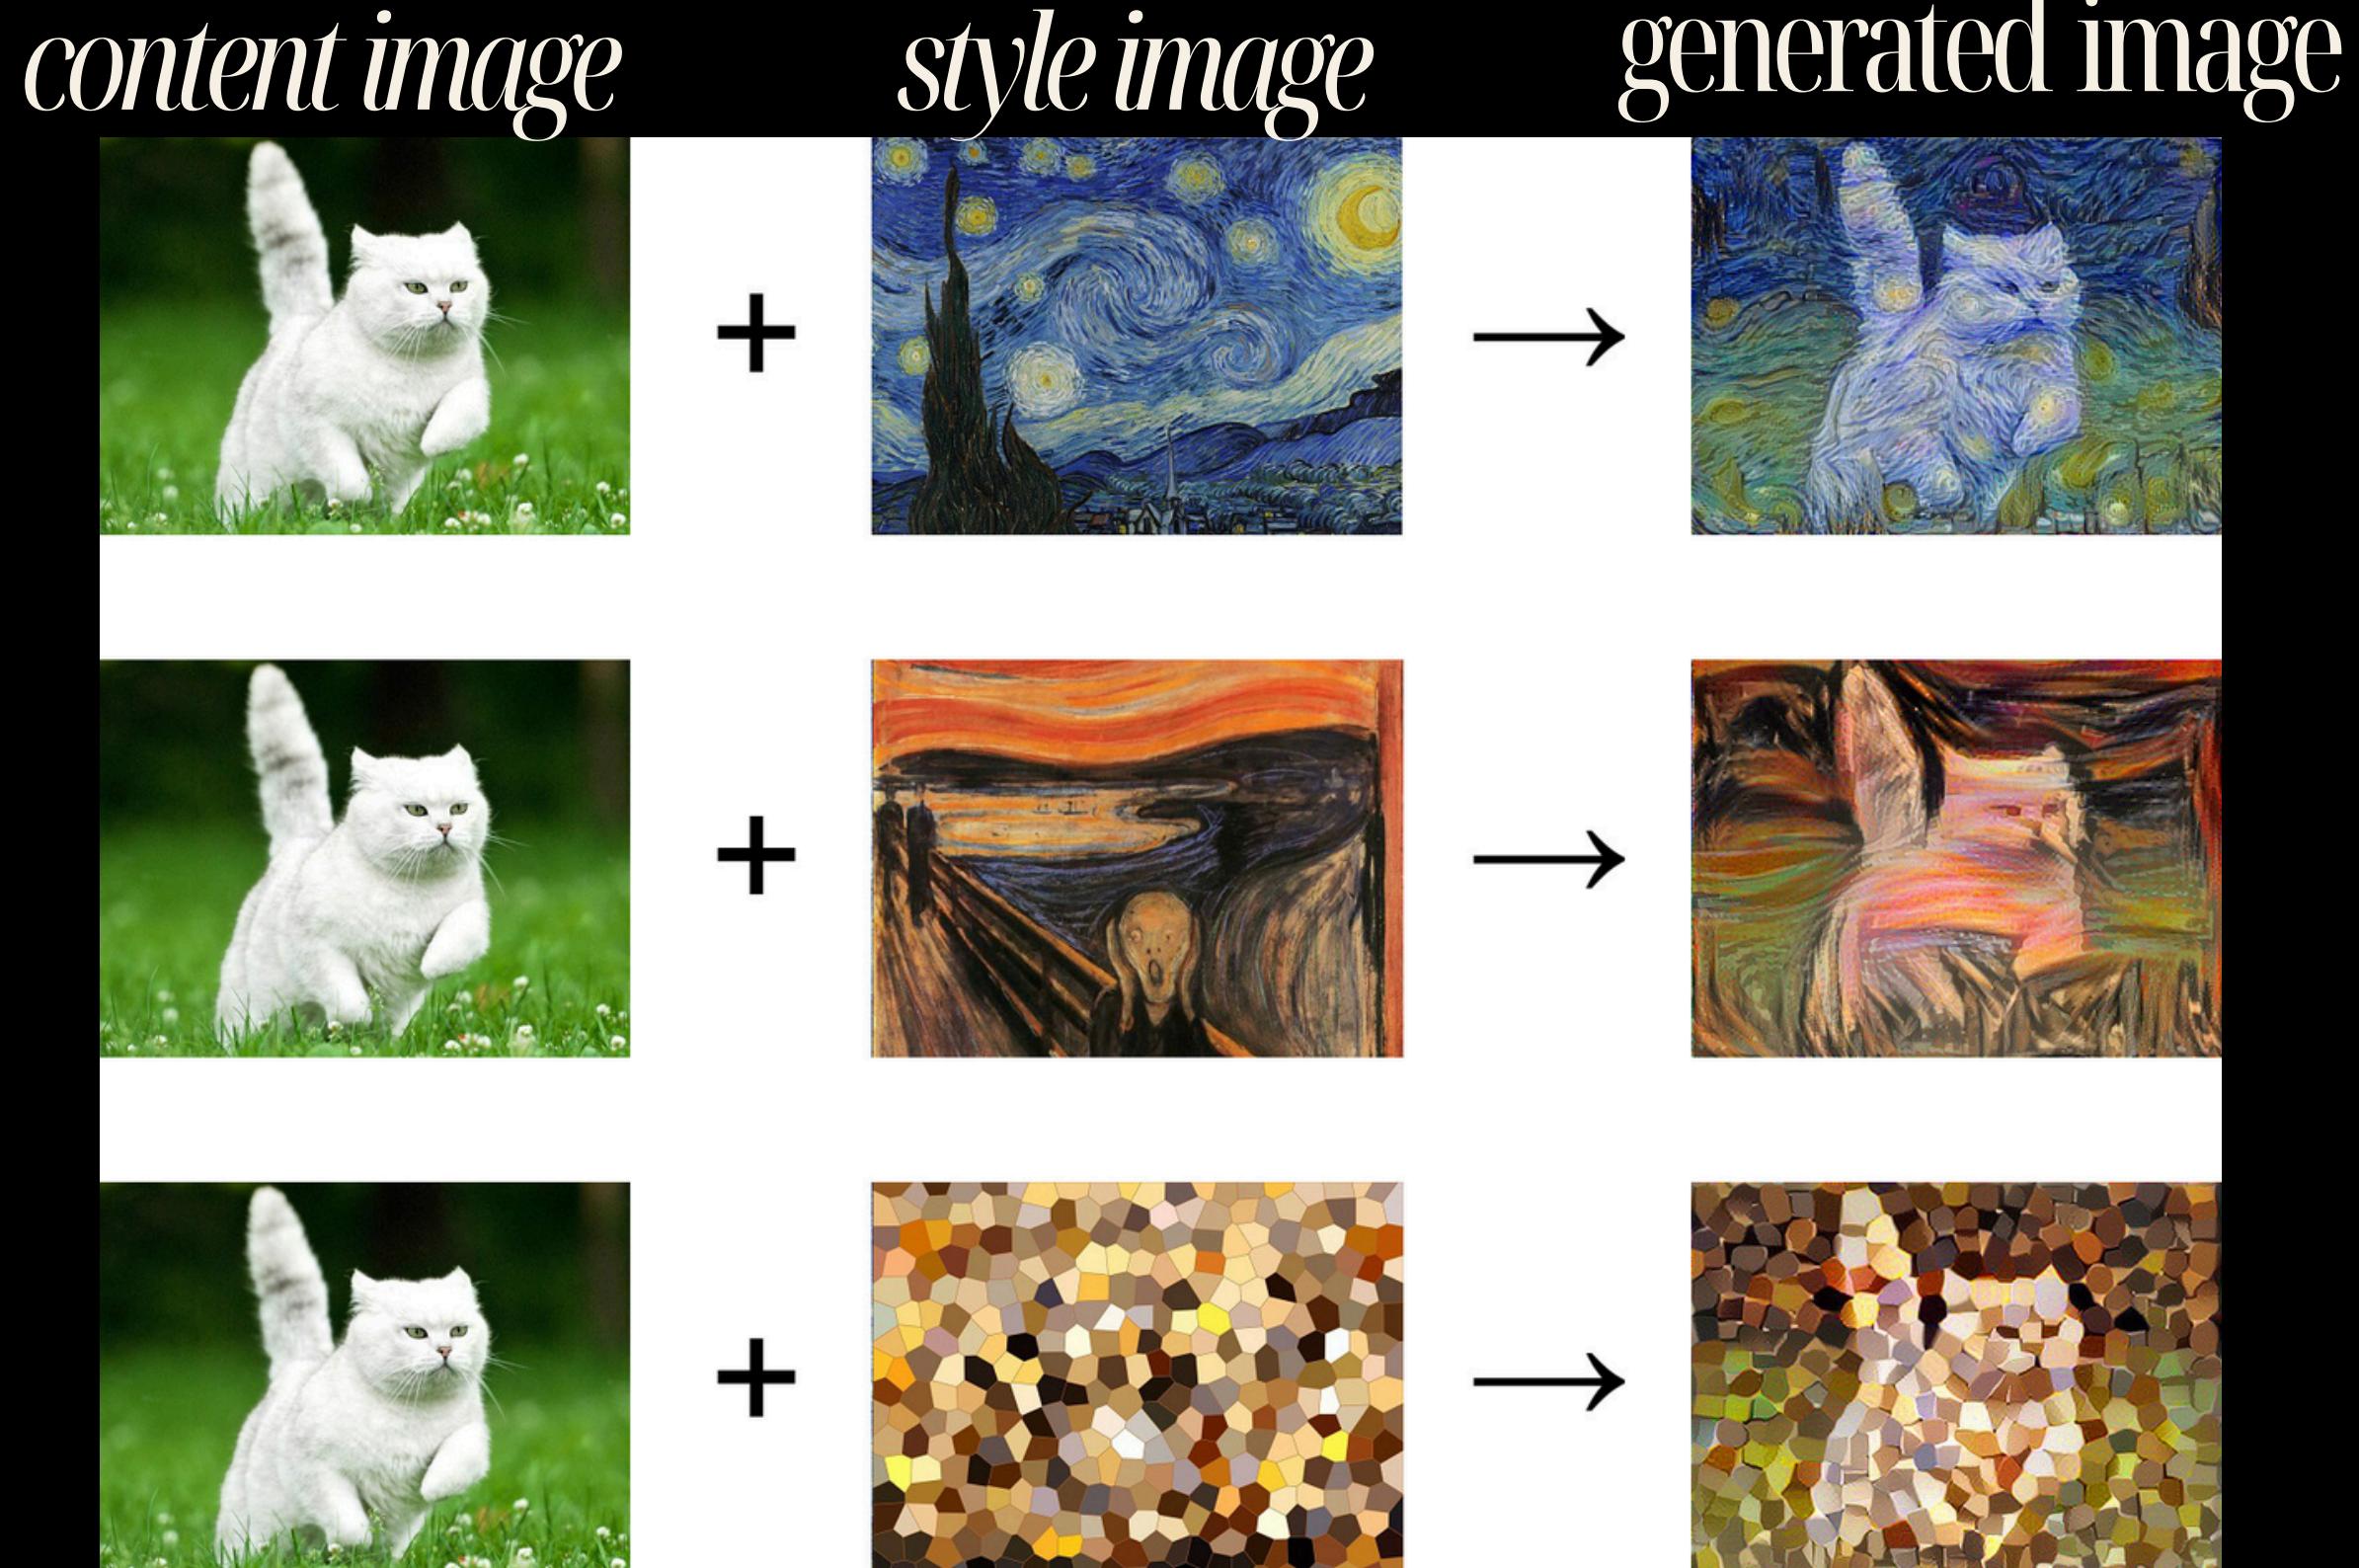
# *Neural Style Transfer*

Under Supervision : Dr. shimaa yosry  
presented by Hana Nabhan

# Fusing Reality with Art

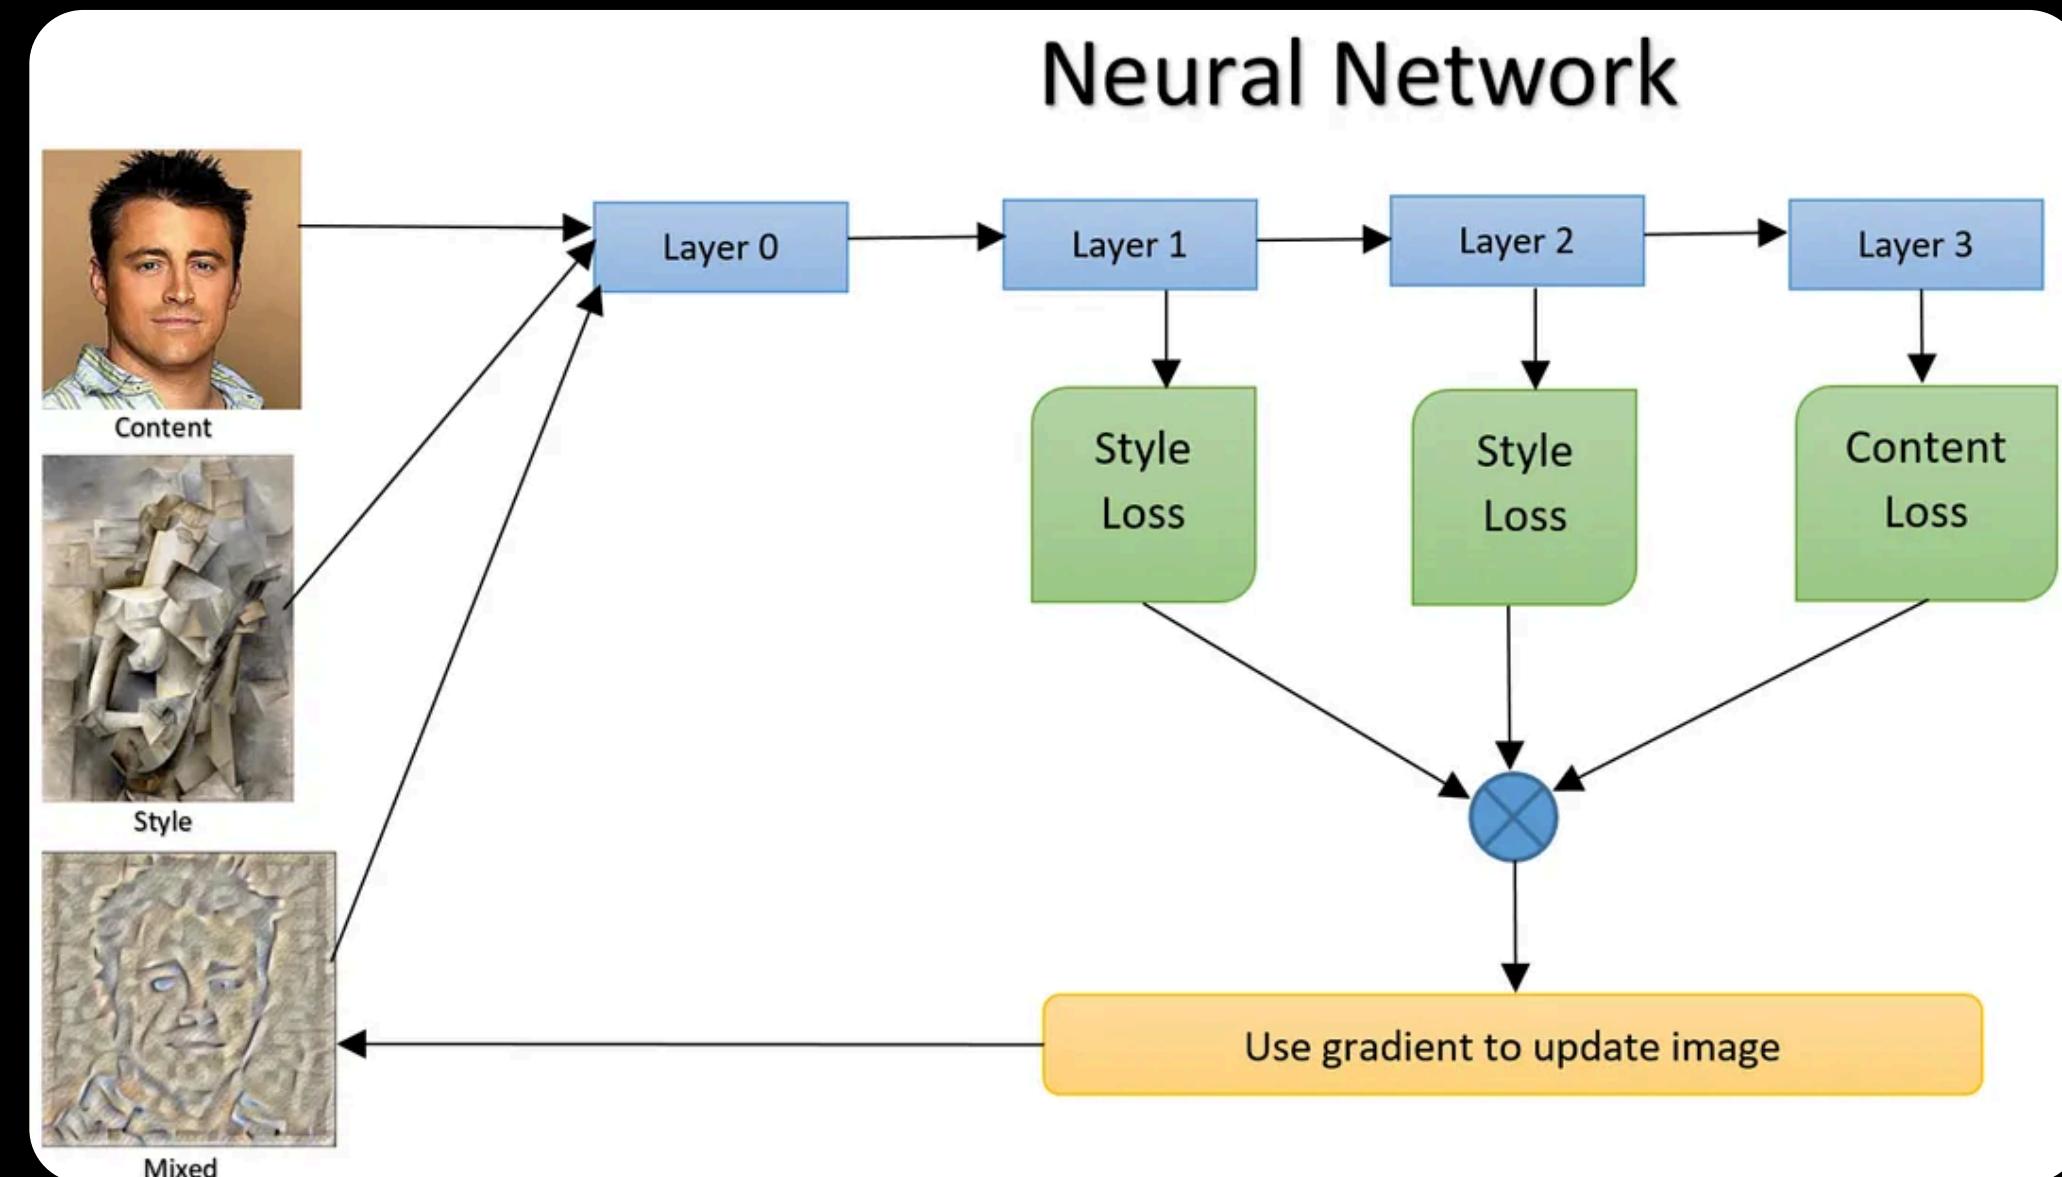
NST is a computer vision technique used to blend three images.

- Photo & Video Editing
- Art & Design
- Gaming & Virtual Reality (VR)
- Data Augmentation
- Interface Customization
- Film & Photography



# *Seeing Like a Machine: The NST Process*

**Define Goals**  
loss function is  
created to  
measure how well  
preserves the  
original's content,  
adopts the new  
style, and remains  
visually coherent.



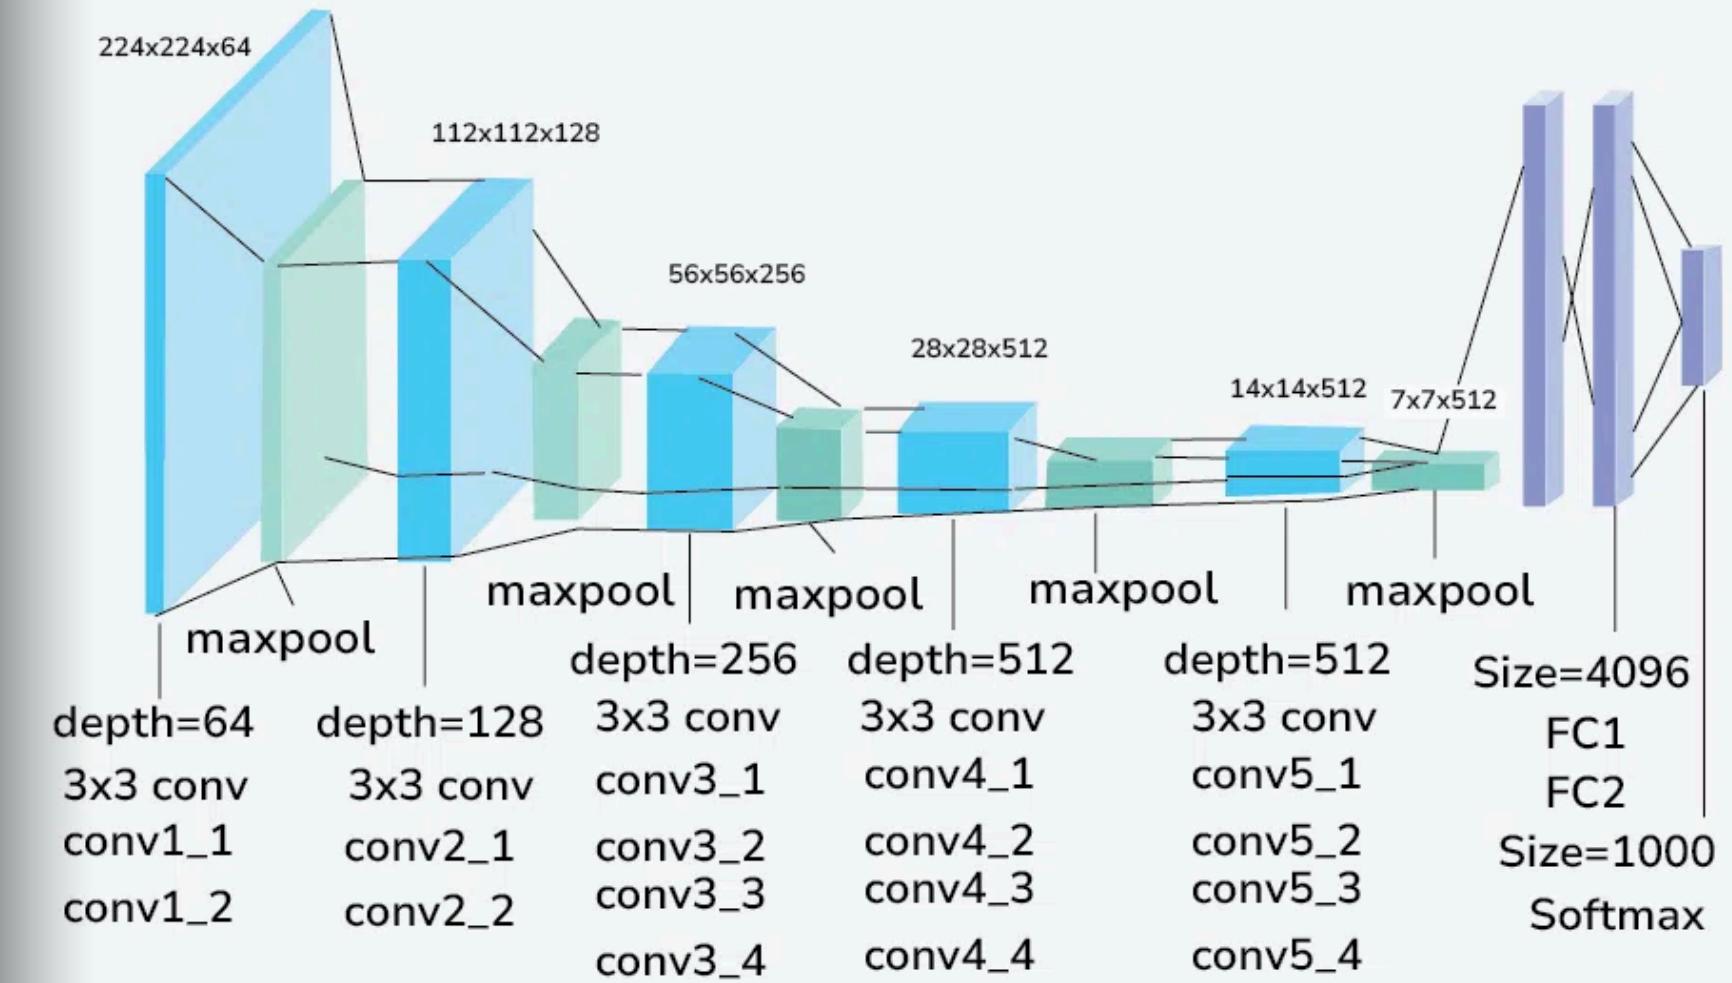
**Perceive**  
(VGG19) is used to  
extract feature

**Optimize**  
adjust to minimize  
the total loss

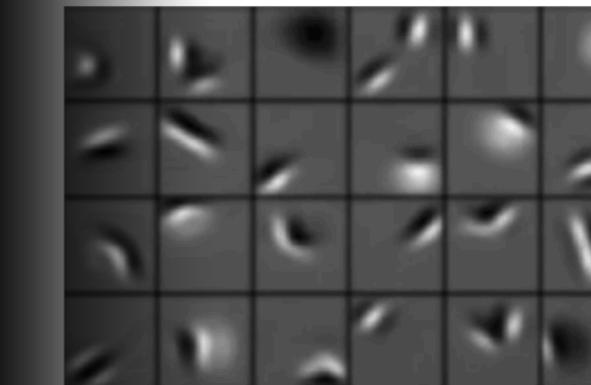
# VGG19 Structure

Trained on ImageNet, VGG19 has learned a rich hierarchy of visual features that can distinguish between high-level structure and low-level texture.

## VGG19 Architecture



Low-level features



Edges, dark spots

Mid-level features



Eyes, ears, nose

High-level features



Facial structure

# The Three Goals of Optimization

## CONTENT LOSS

---

"Keep the photo's structure."  
Measures the squared-error (L<sub>2</sub> distance) between the feature representations of the content image and the generated image.

## STYLE LOSS

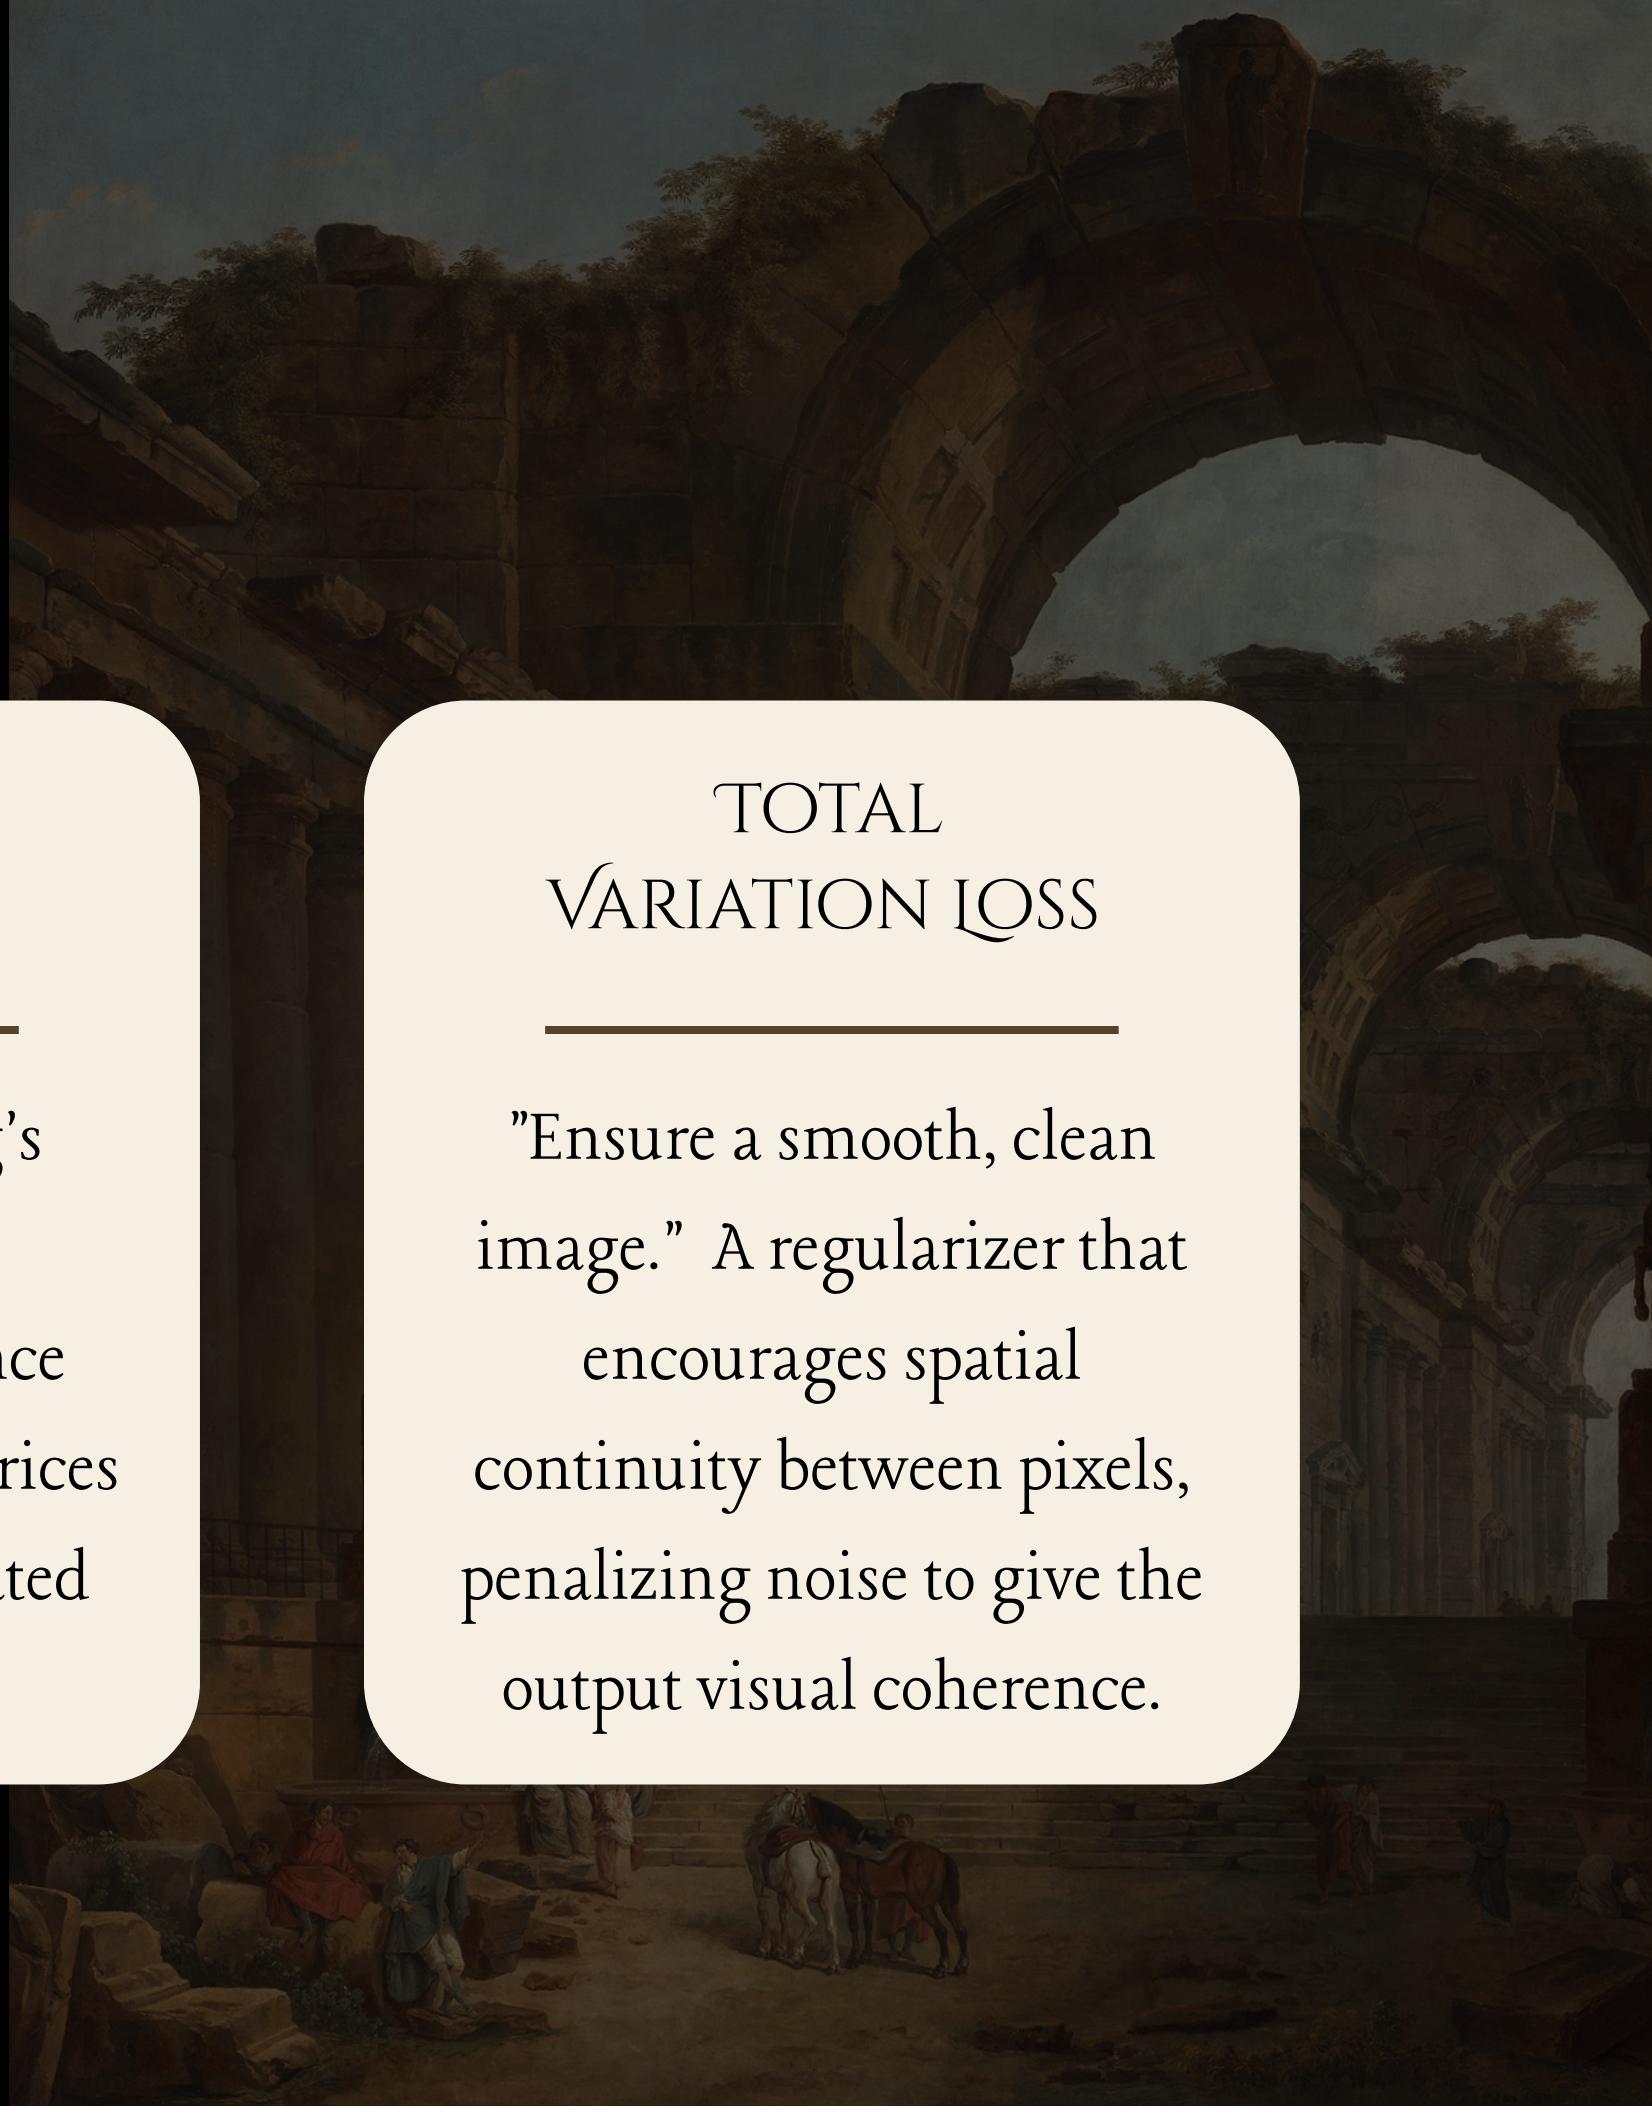
---

"Mimic the painting's texture."  
Measures the difference between the Gram matrices of the style and generated images.

## TOTAL VARIATION LOSS

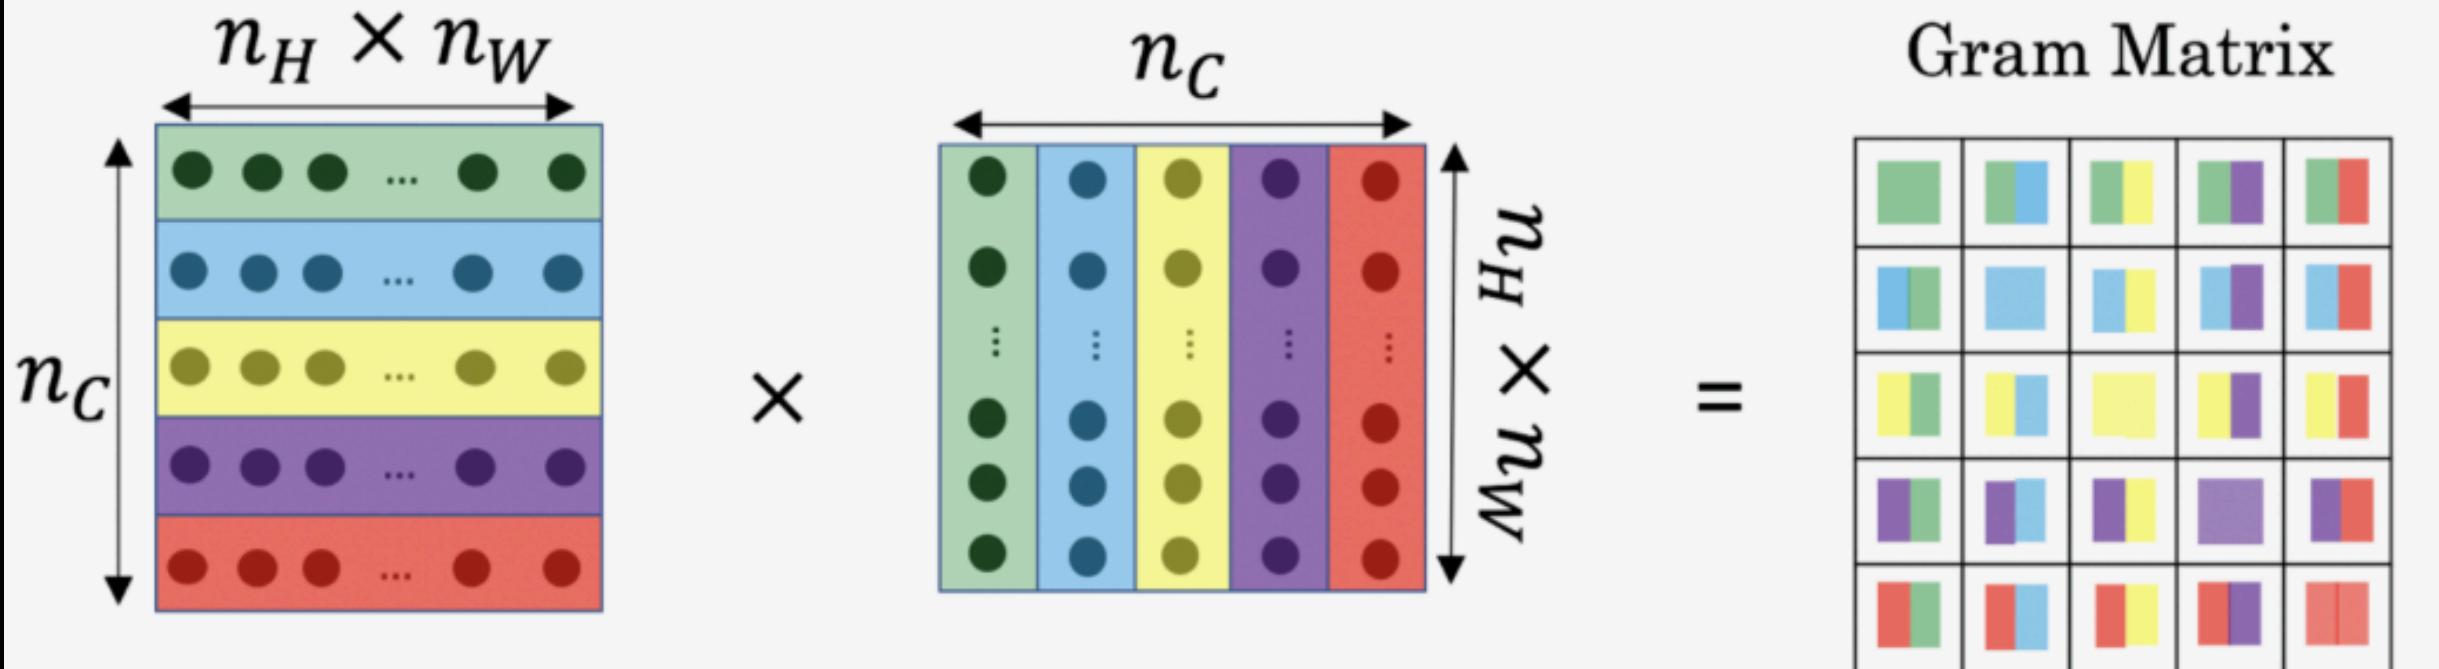
---

"Ensure a smooth, clean image." A regularizer that encourages spatial continuity between pixels, penalizing noise to give the output visual coherence.



# The Gram Matrix

$$G_{ij}^l = \sum F_{ik}^l F_{jk}^l.$$



Gram Matrix

green	green	green	green	green
blue	blue	blue	blue	blue
yellow	yellow	yellow	yellow	yellow
purple	purple	purple	purple	purple
red	red	red	red	red

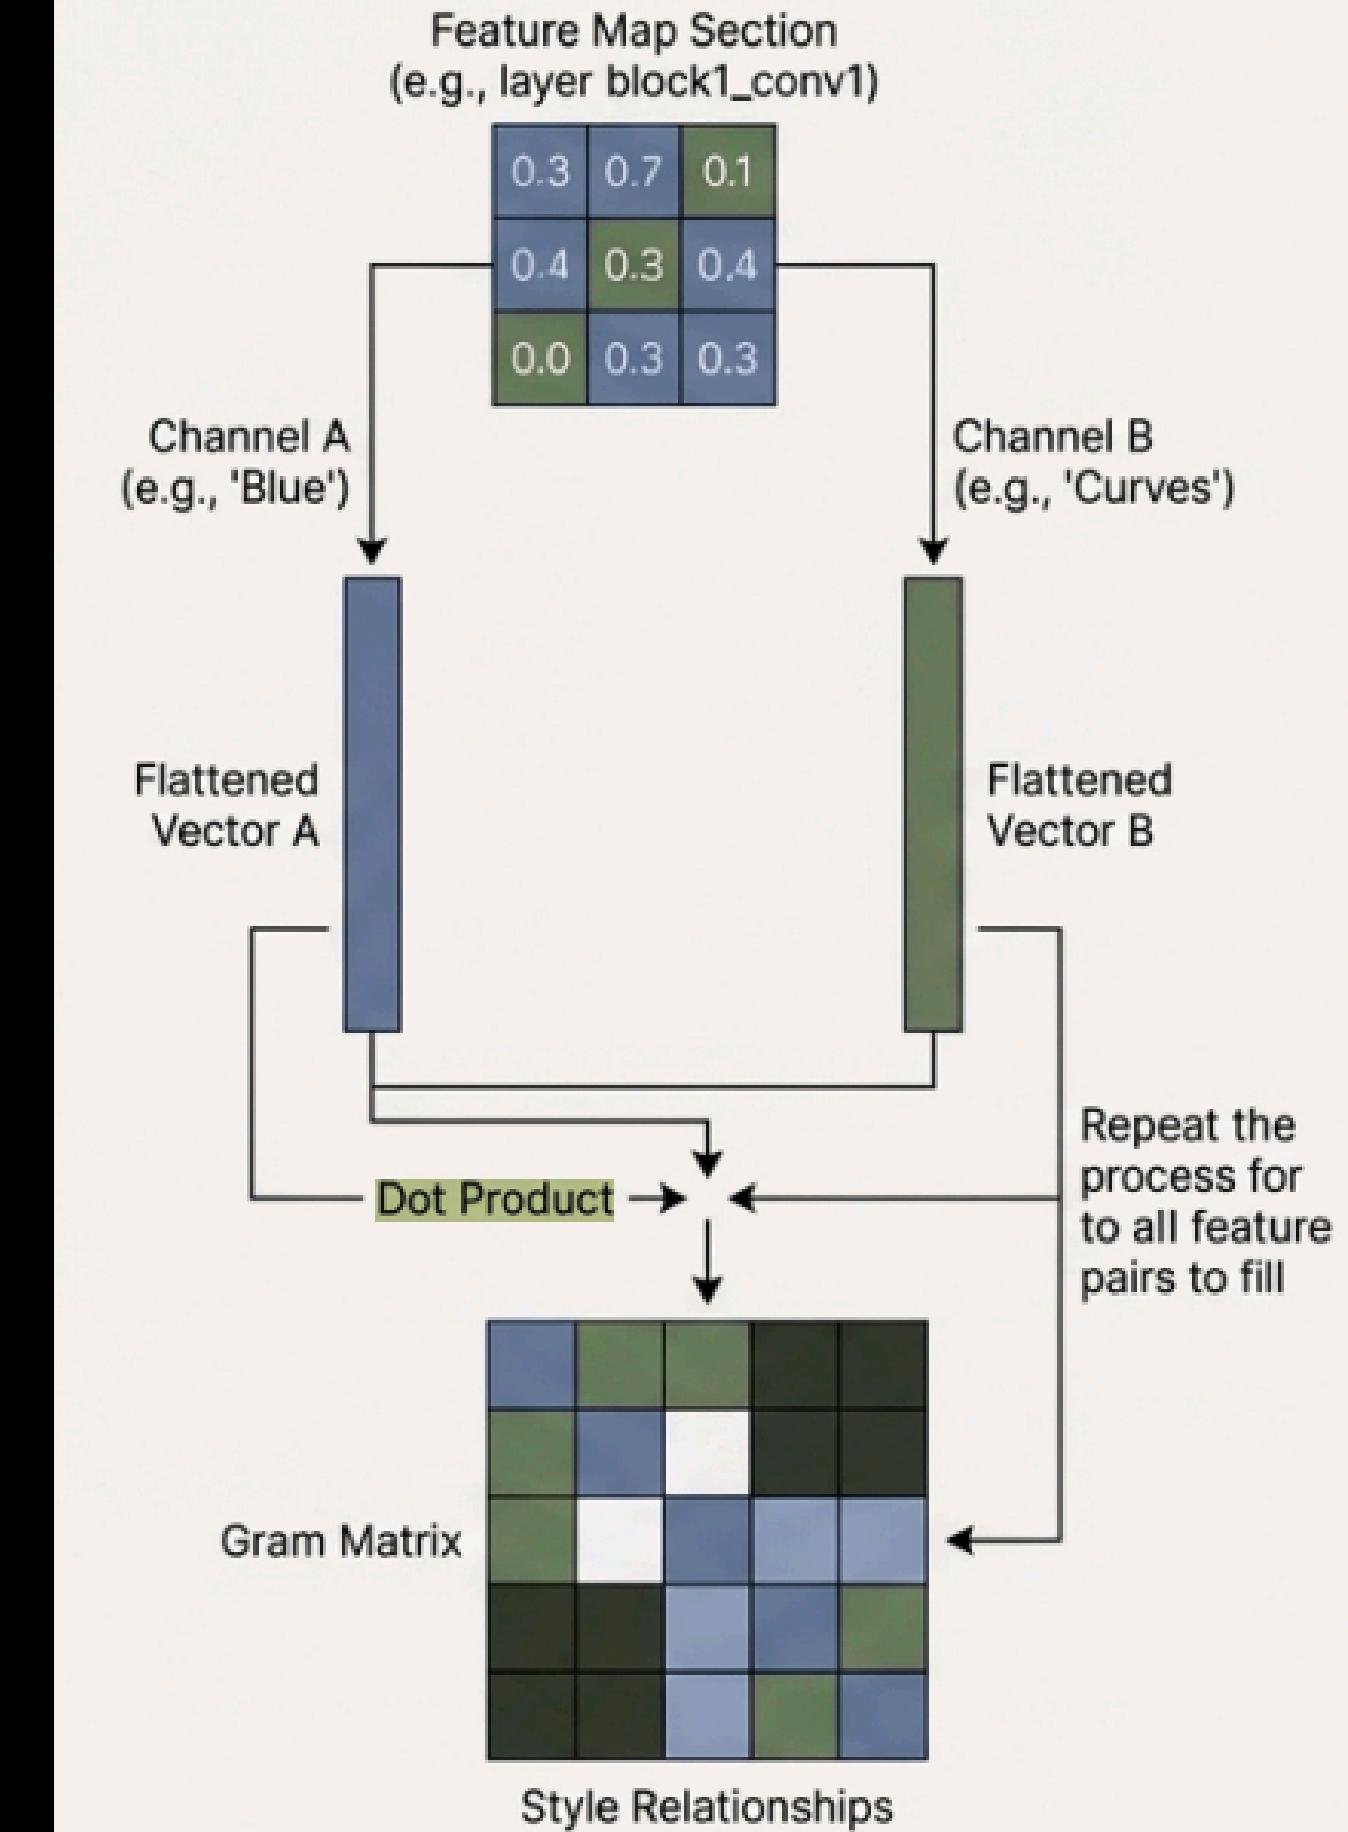
Correlation between filter

## What it is?

The Gram Matrix measures the correlations between different filter responses in a single layer. It captures which features tend to activate together, representing the image's texture and patterns, independent of their specific location.

# The Gram Matrix

**The Core Idea Dot Product "In the original painting, whenever I see 'Blue', I also see 'Curved Lines'. Does your image do that too? This relationship (Blue + Curves) is calculated using the Gram Matrix."**



# *Loss functions*

$$\mathcal{L}_{content}(\vec{p}, \vec{x}, l) = \frac{1}{2} \sum_{i,j} (F_{ij}^l - P_{ij}^l)^2$$

$$E_l = \frac{1}{4N_l^2 M_l^2} \sum_{i,j} (G_{ij}^l - A_{ij}^l)^2$$

$$\mathcal{L}_{style}(\vec{a}, \vec{x}) = \sum_{l=0}^L w_l E_l$$

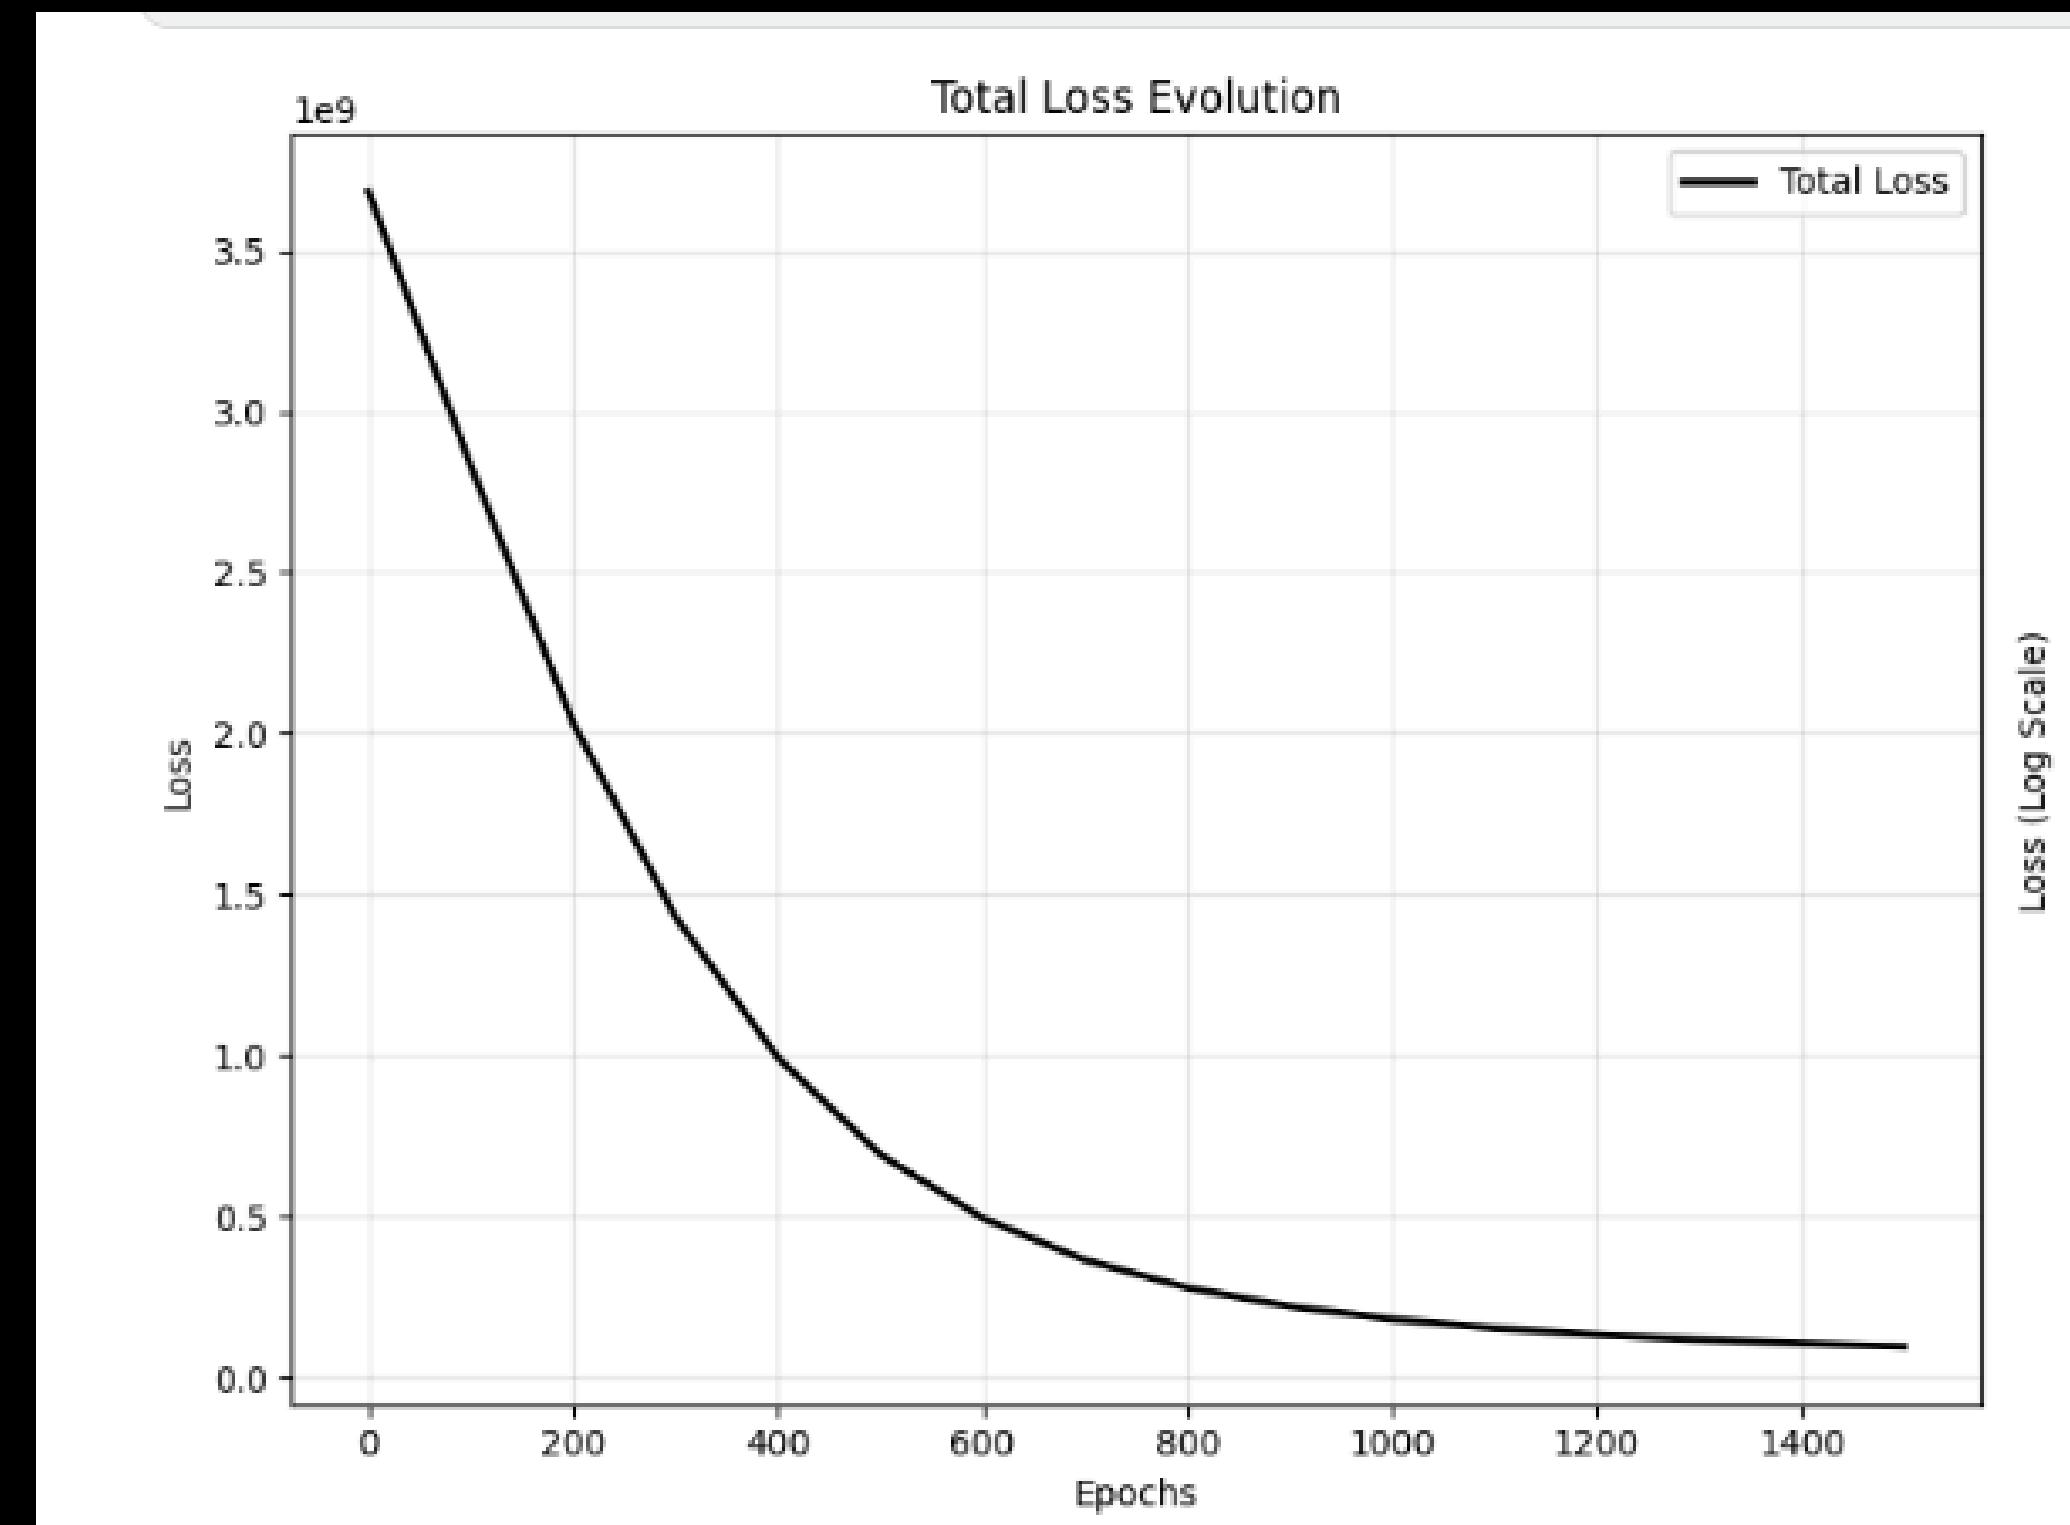
$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}_{content}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

Adam is a gradient-based optimizer commonly used in deep learning. The process uses a `@tf.function` called `train\_step` that calculates the total loss and applies gradients in each step.

## Hyperparameters

content\_weight = 1000.0, style\_weight = 0.1,  
total\_variation\_weight = 30.

*adam*

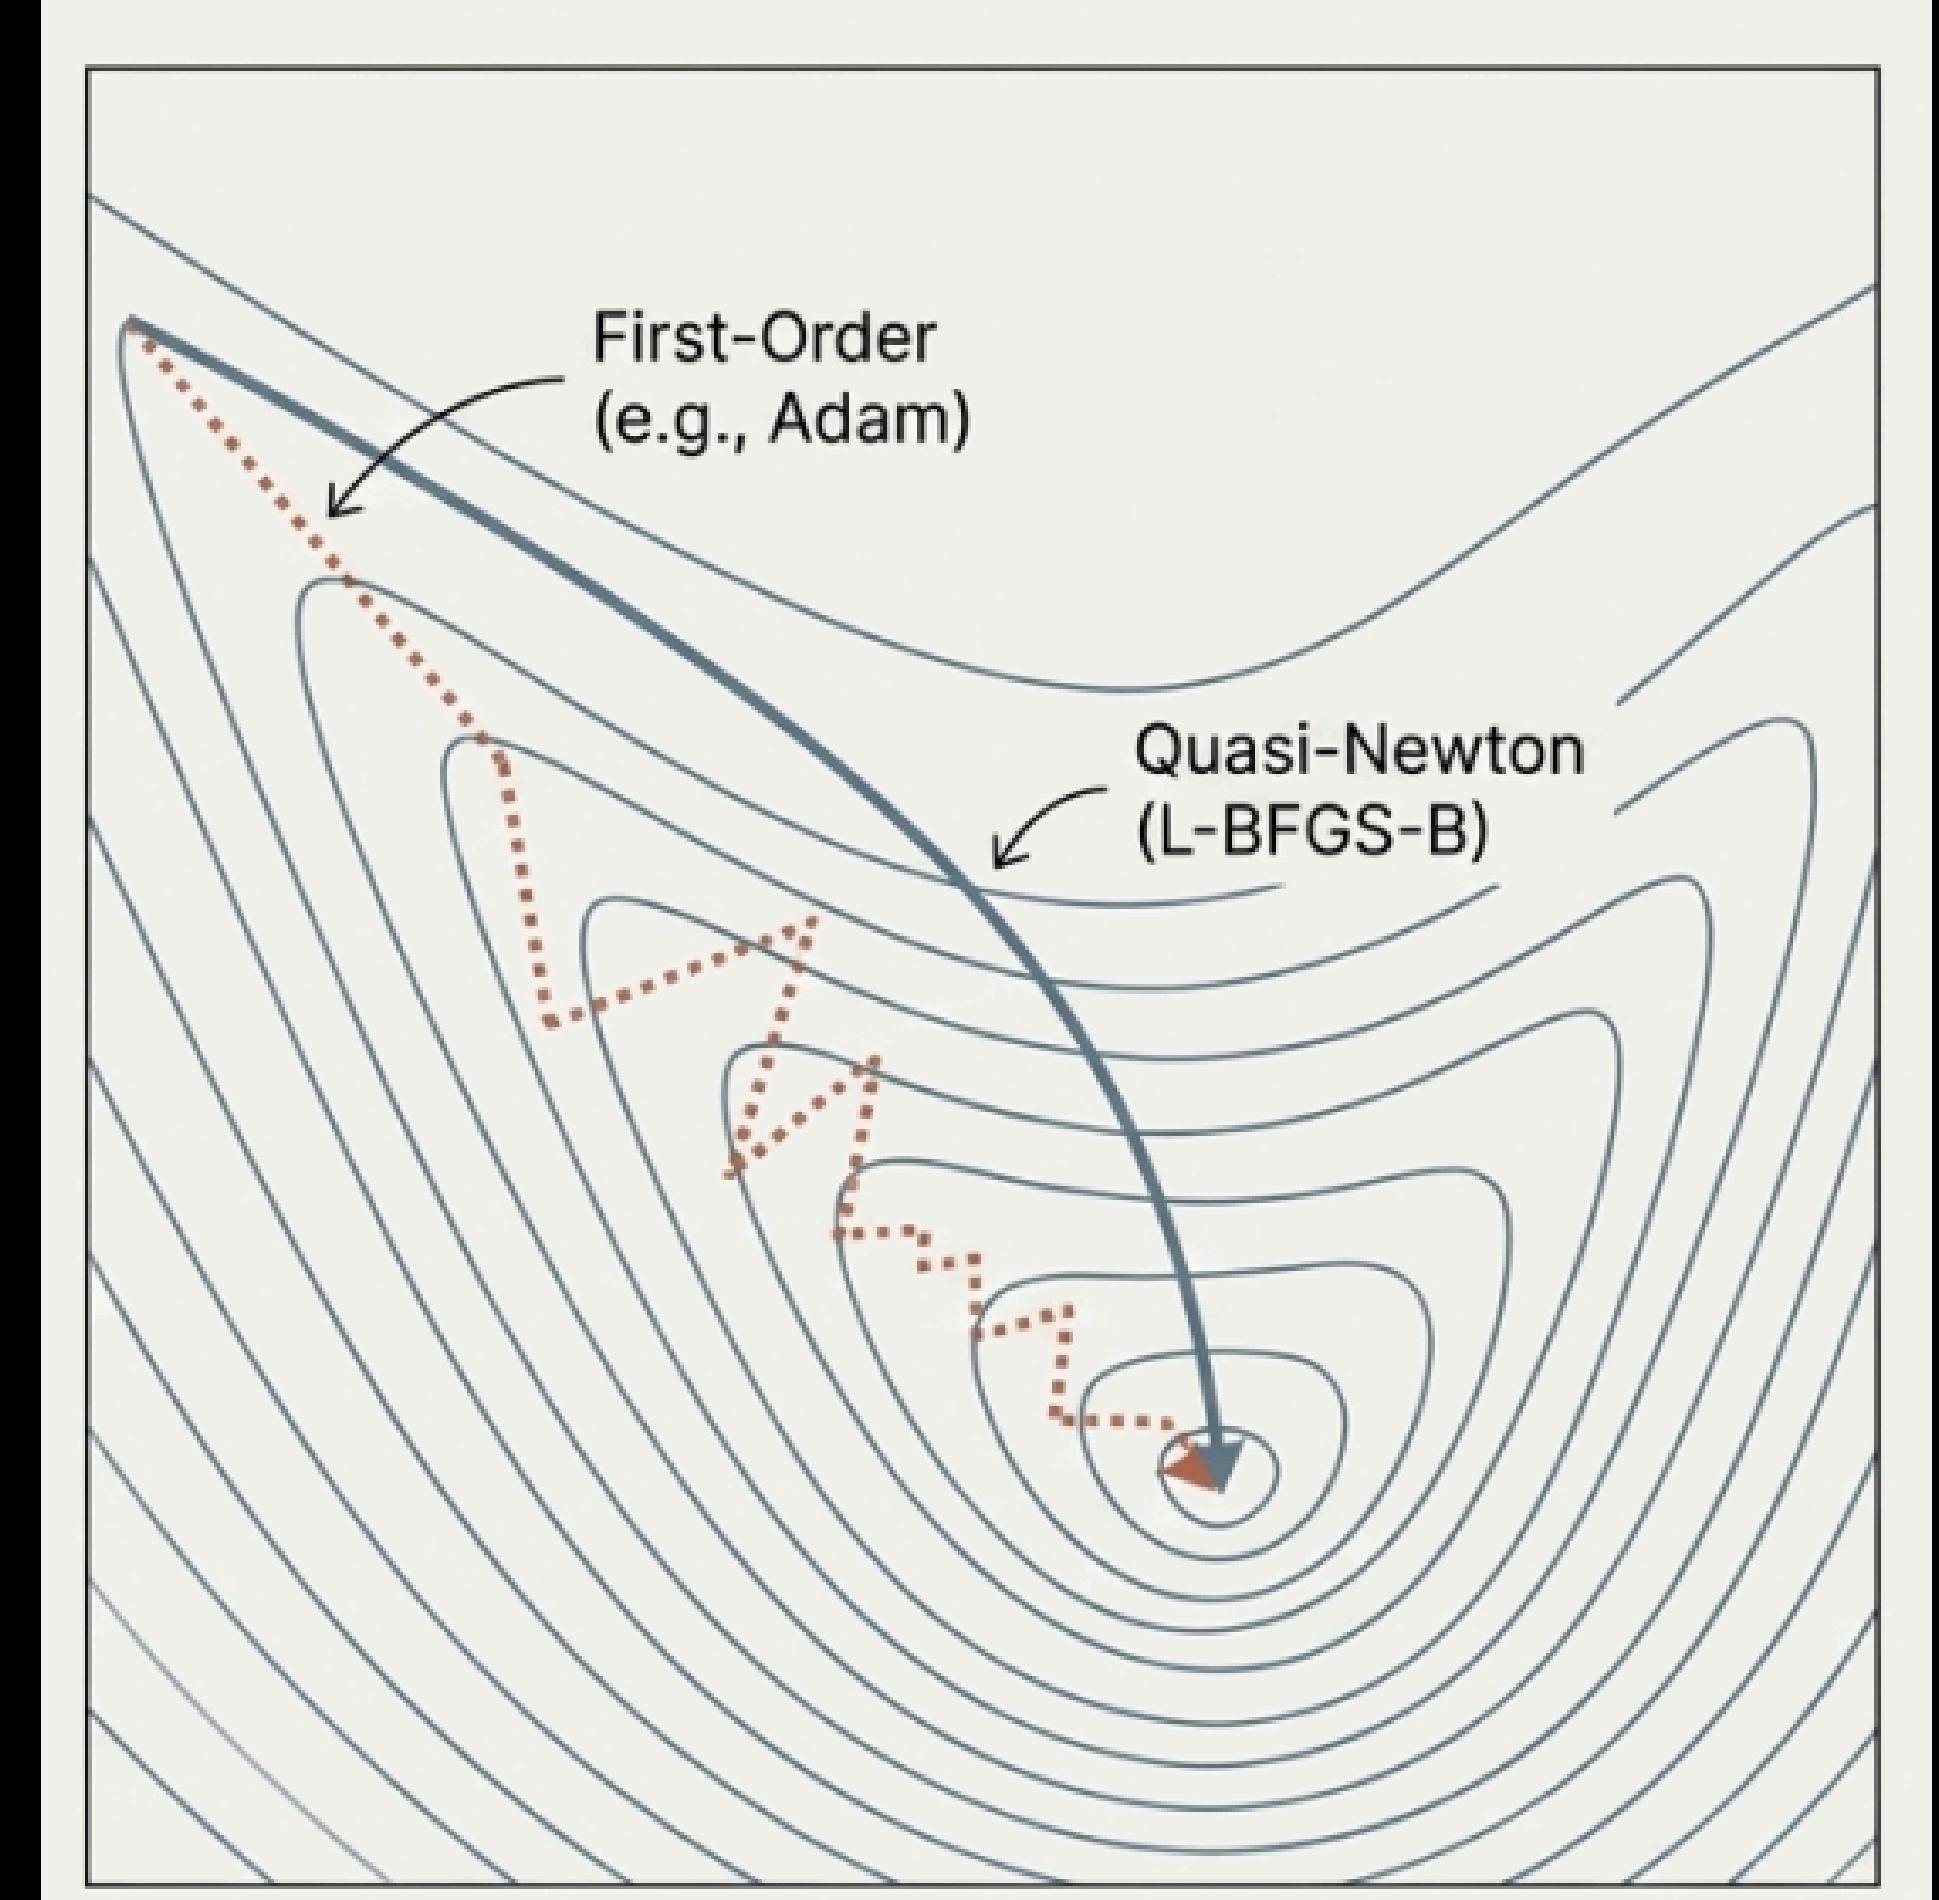


# *L-BFGS-B*

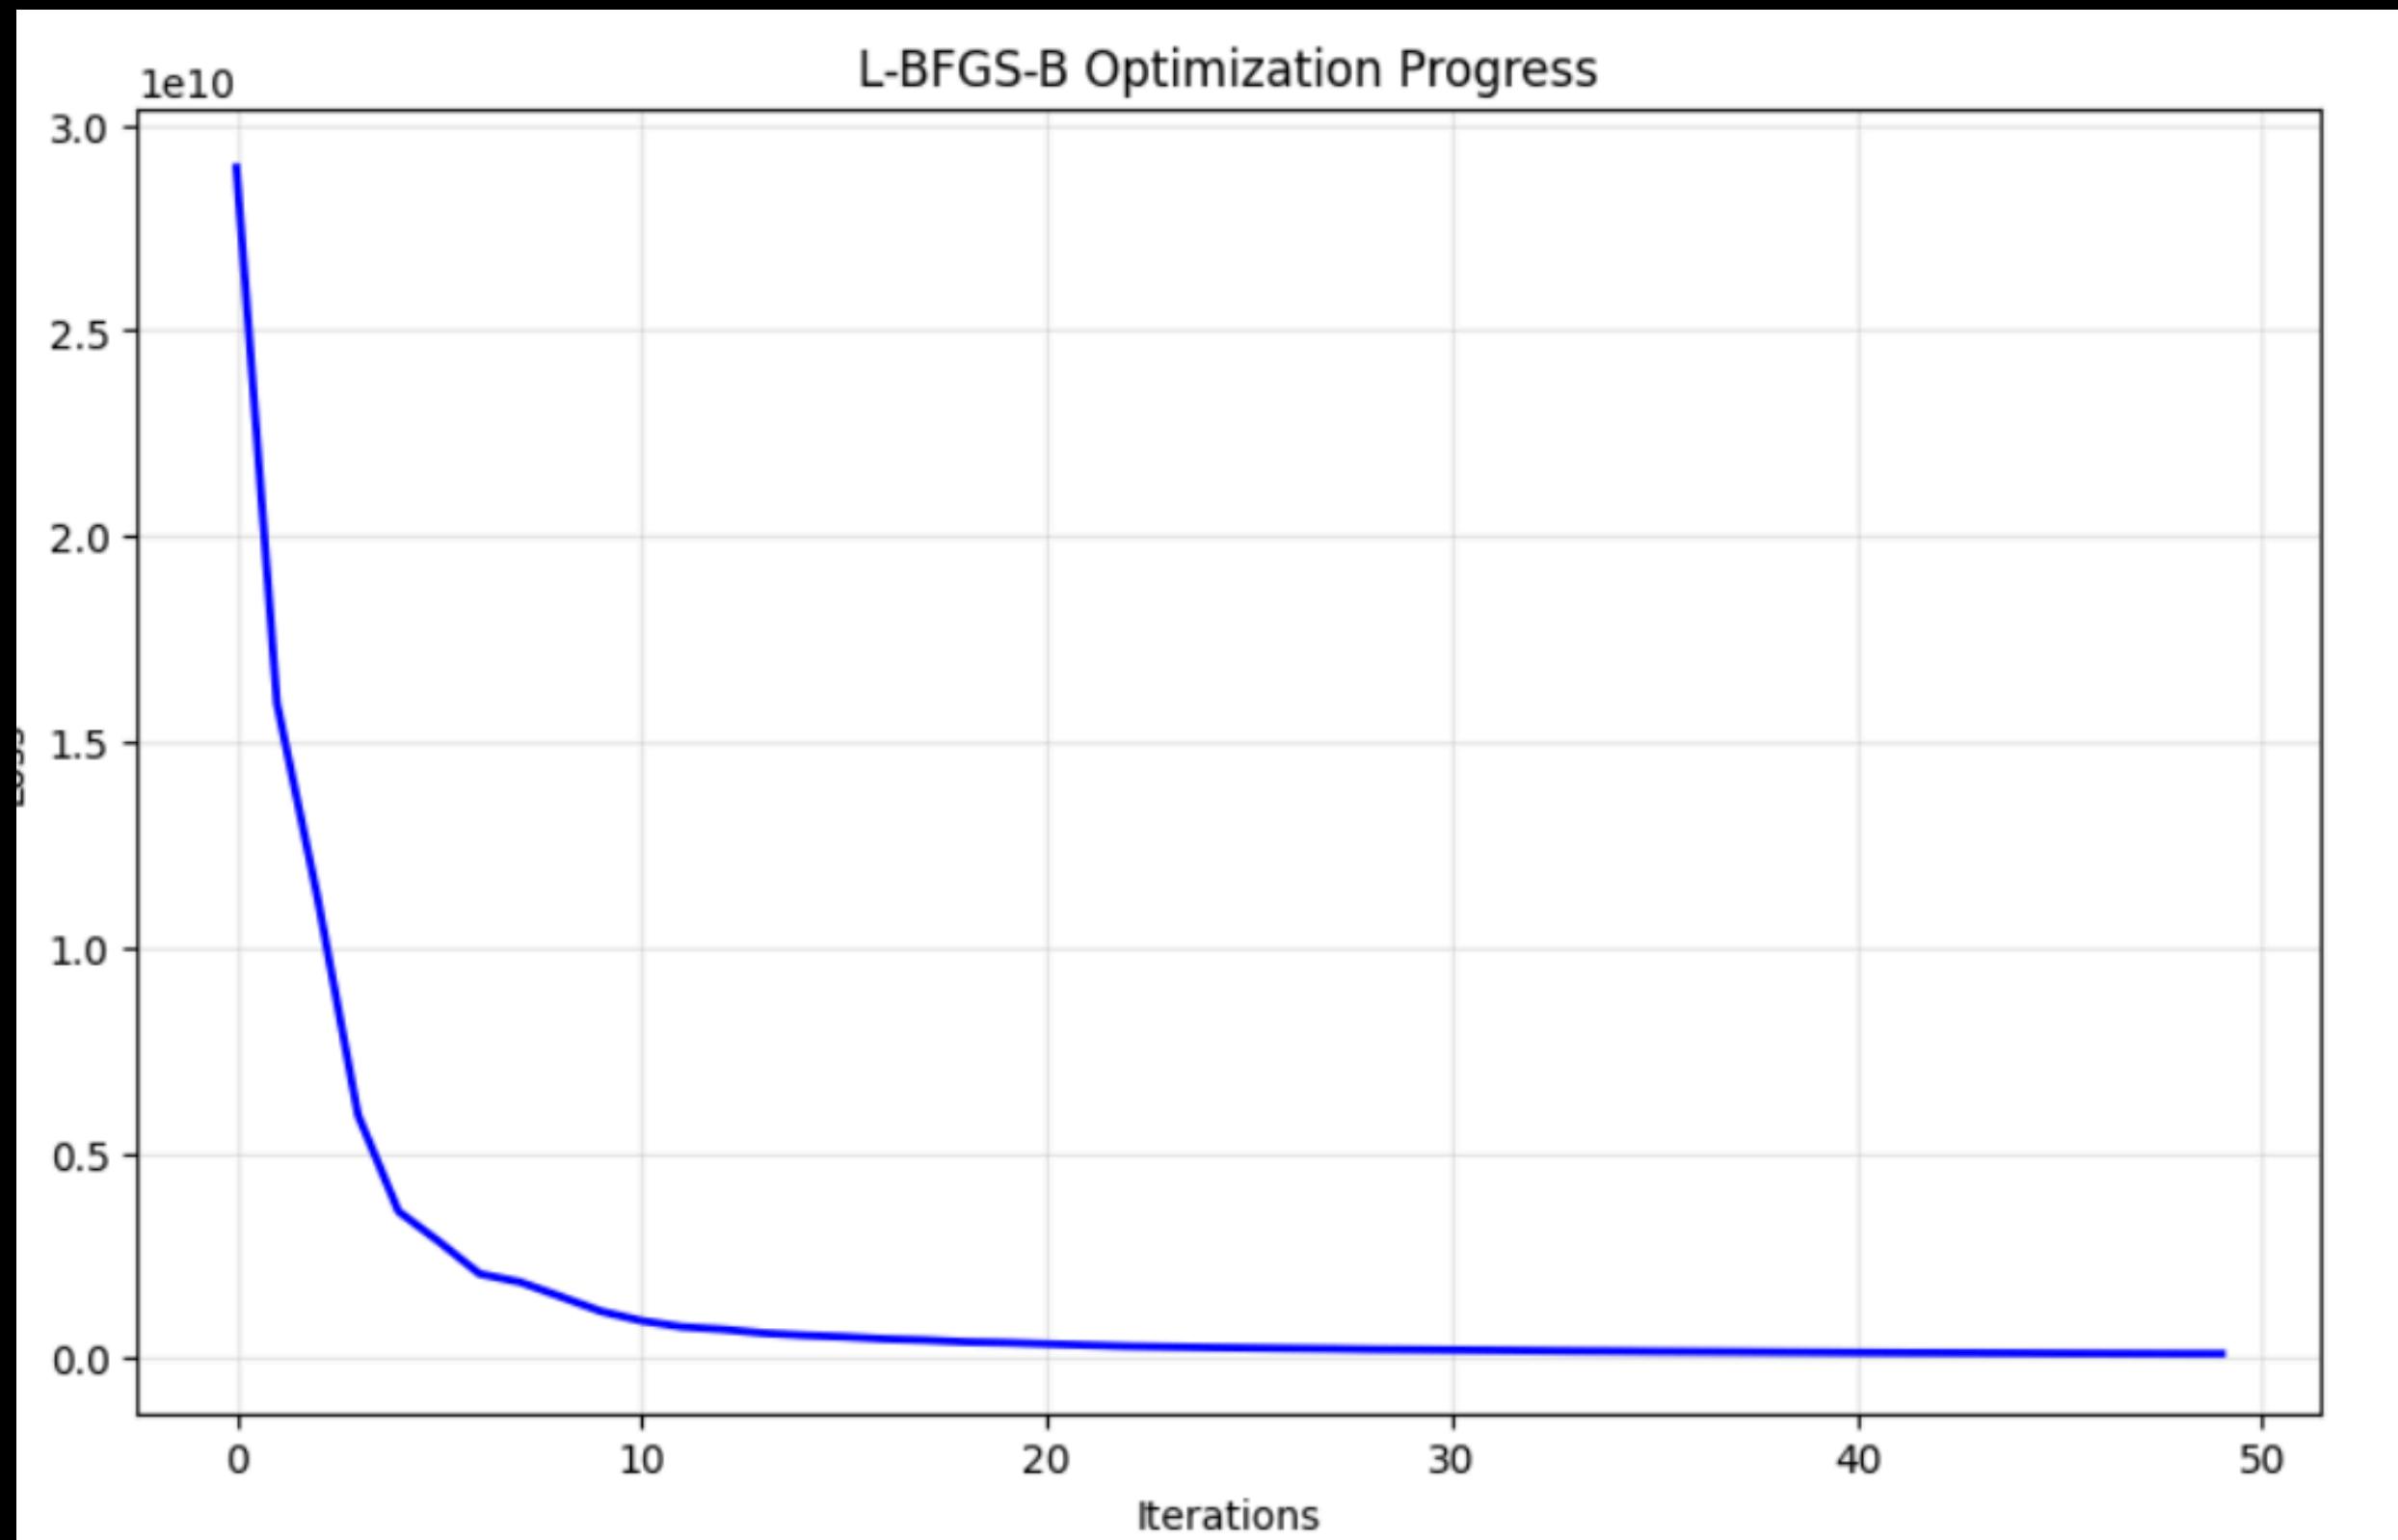
Why L-BFGS-B?

It is a quasi-Newton method that can converge faster and produce sharper results.

It uses smarter mathematics by estimating the curvature of the loss landscape, not just the slope.

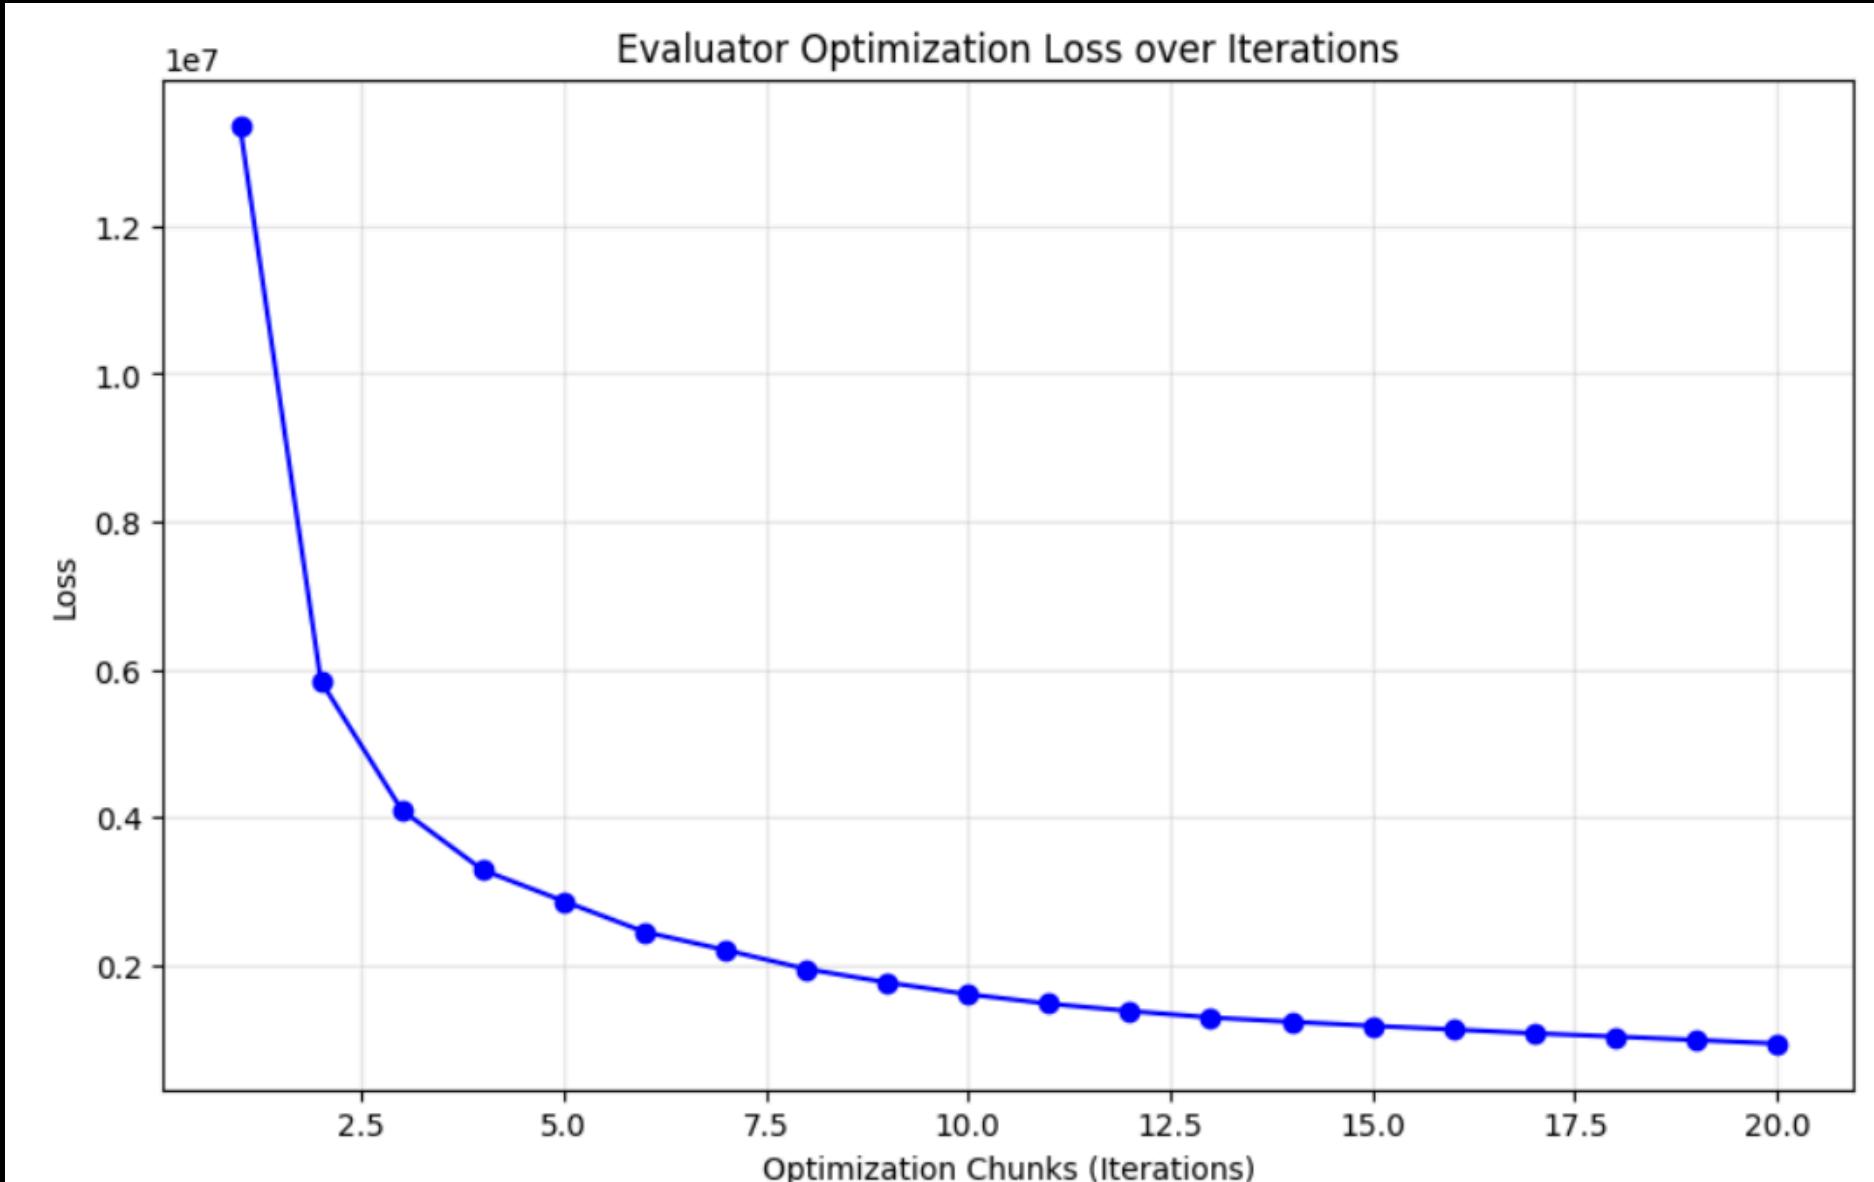


## L-BFGS-B Optimization Progress



*L-BFGS-B*

# *A New Strategy and a Hidden Cost (Evaluator)*



The ``scipy.optimize`` function requests loss and gradients in separate calls. A naive implementation calculates the full model pass twice-once for the loss and again for the gradients, which doubles the runtime.

### The Problem:

Redundant computations because the optimizer asks for loss and gradients in separate function calls.

**The Solution:** An `Evaluator` class computes both loss and gradients in a single pass. It caches these results, so subsequent calls for either value within the same optimization step are instantaneous.

### The Solution:

An `Evaluator` class computes both loss and gradients in a single pass. It caches these results, so subsequent calls for either value within the same optimization step are instantaneous.

# *Maximum Efficiency: The Evaluator Pattern*

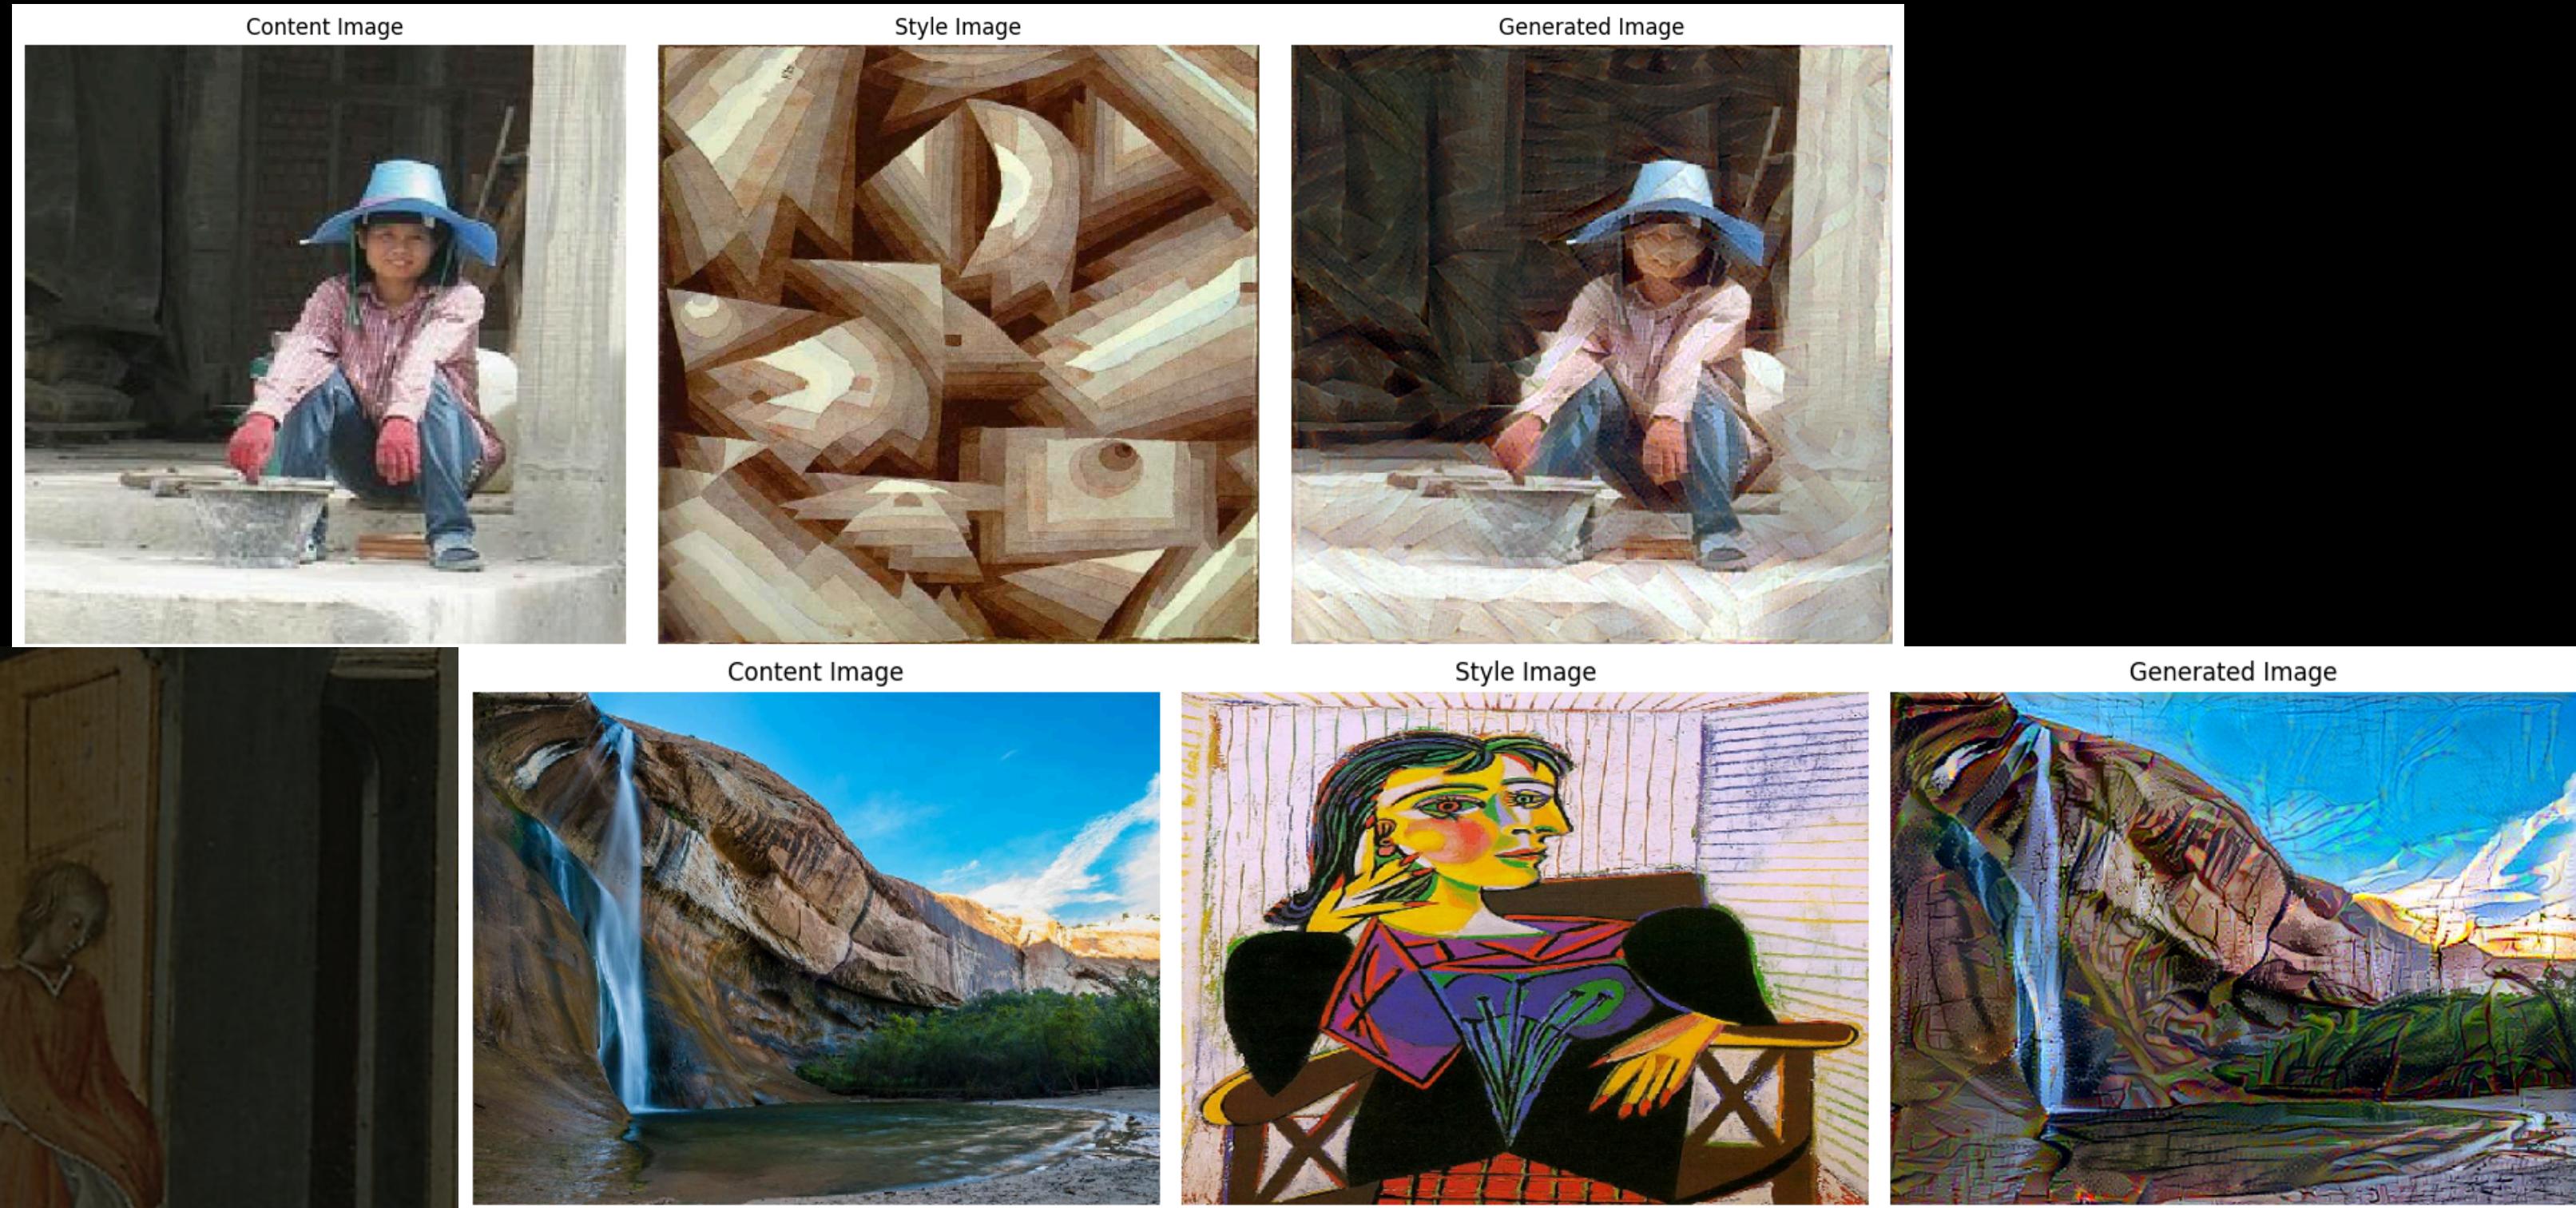
# Performance Review

## Adam vs. L-BFGS

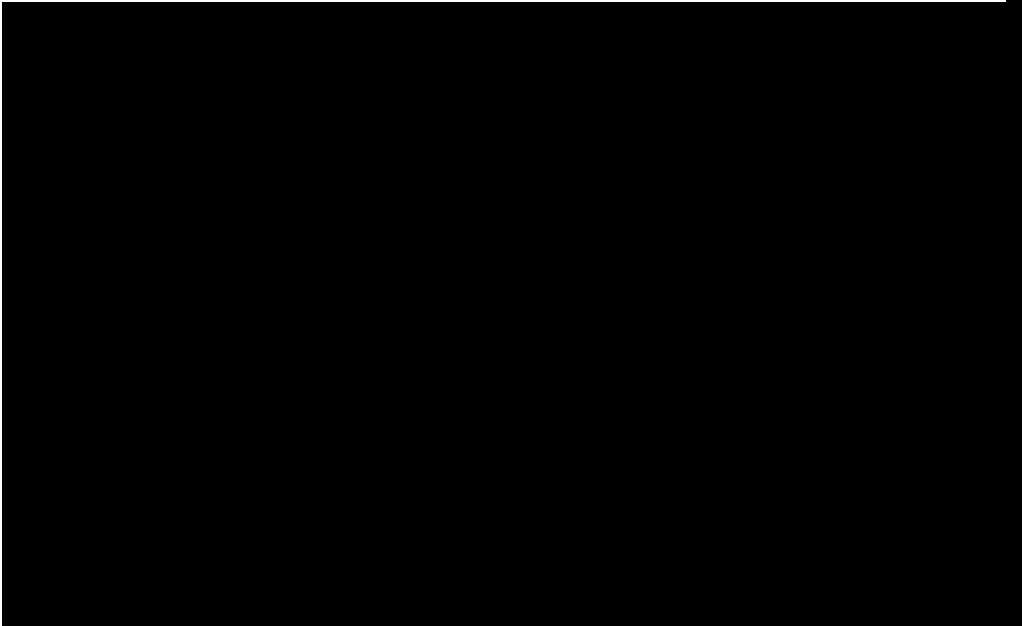
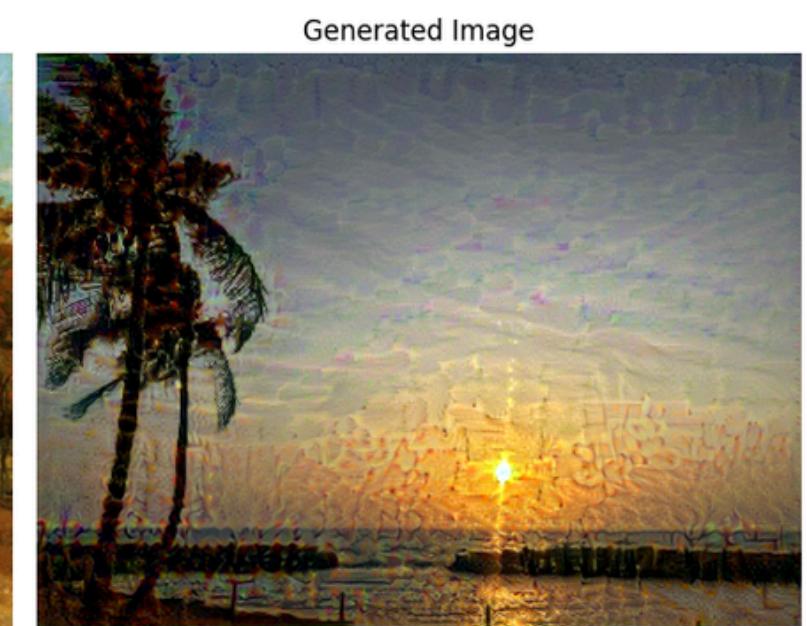
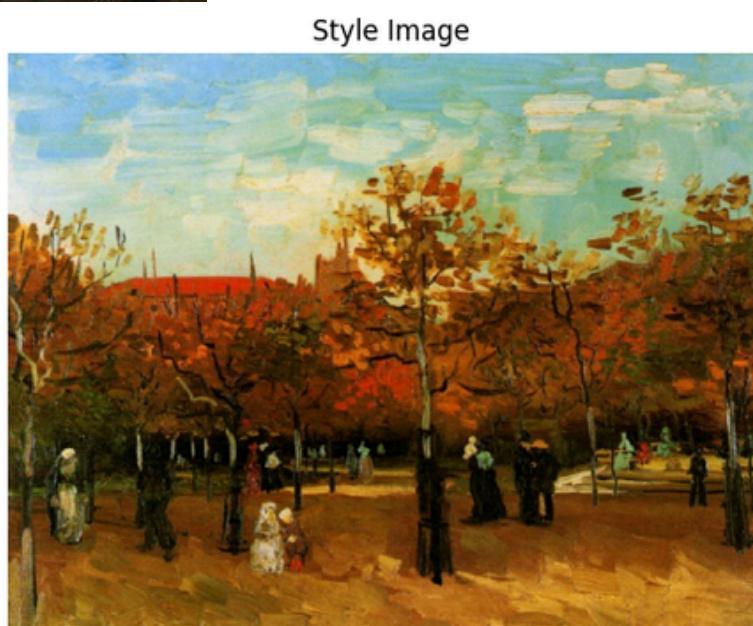
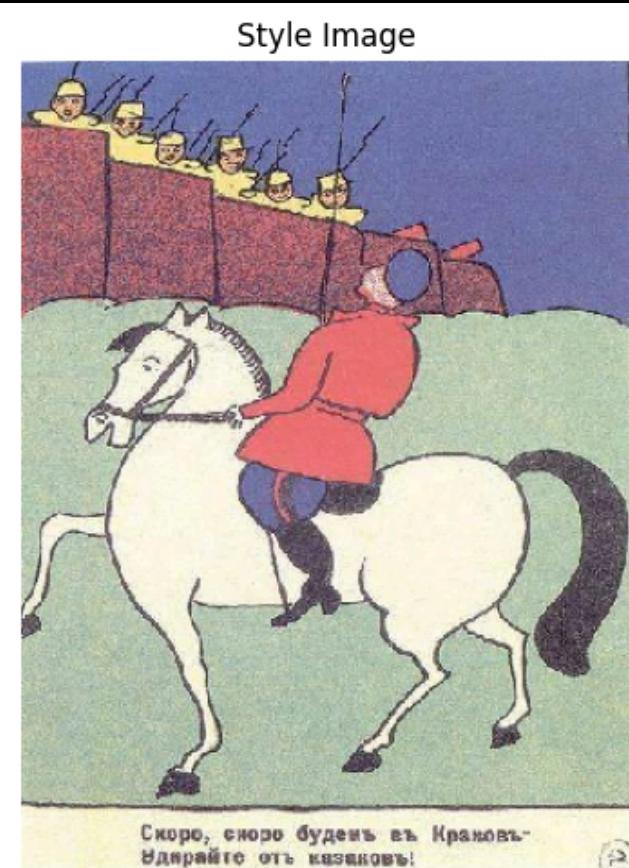
L-BFGS HAS A "HEAVIER" COMPUTATIONAL COST PER STEP, BUT BECAUSE IT FINDS A MORE DIRECT PATH TO THE SOLUTION, THE TOTAL TIME TO REACH A HIGH-QUALITY RESULT IS MUCH FASTER.

Metric	Adam	L-BFGS-B (with Evaluator)
Mechanism	First-Order (Gradient Descent)	Quasi-Newton (Approximates Curvature)
Iterations	Many	Few
Quality	Good	Often Sharper, Higher Fidelity
Implementation	Simpler, integrated into TF	More complex; requires Evaluator pattern
Best For	Robust baseline, easy to implement	Faster convergence to high quality, sharper results

# *The Versatility of the Algorithm*

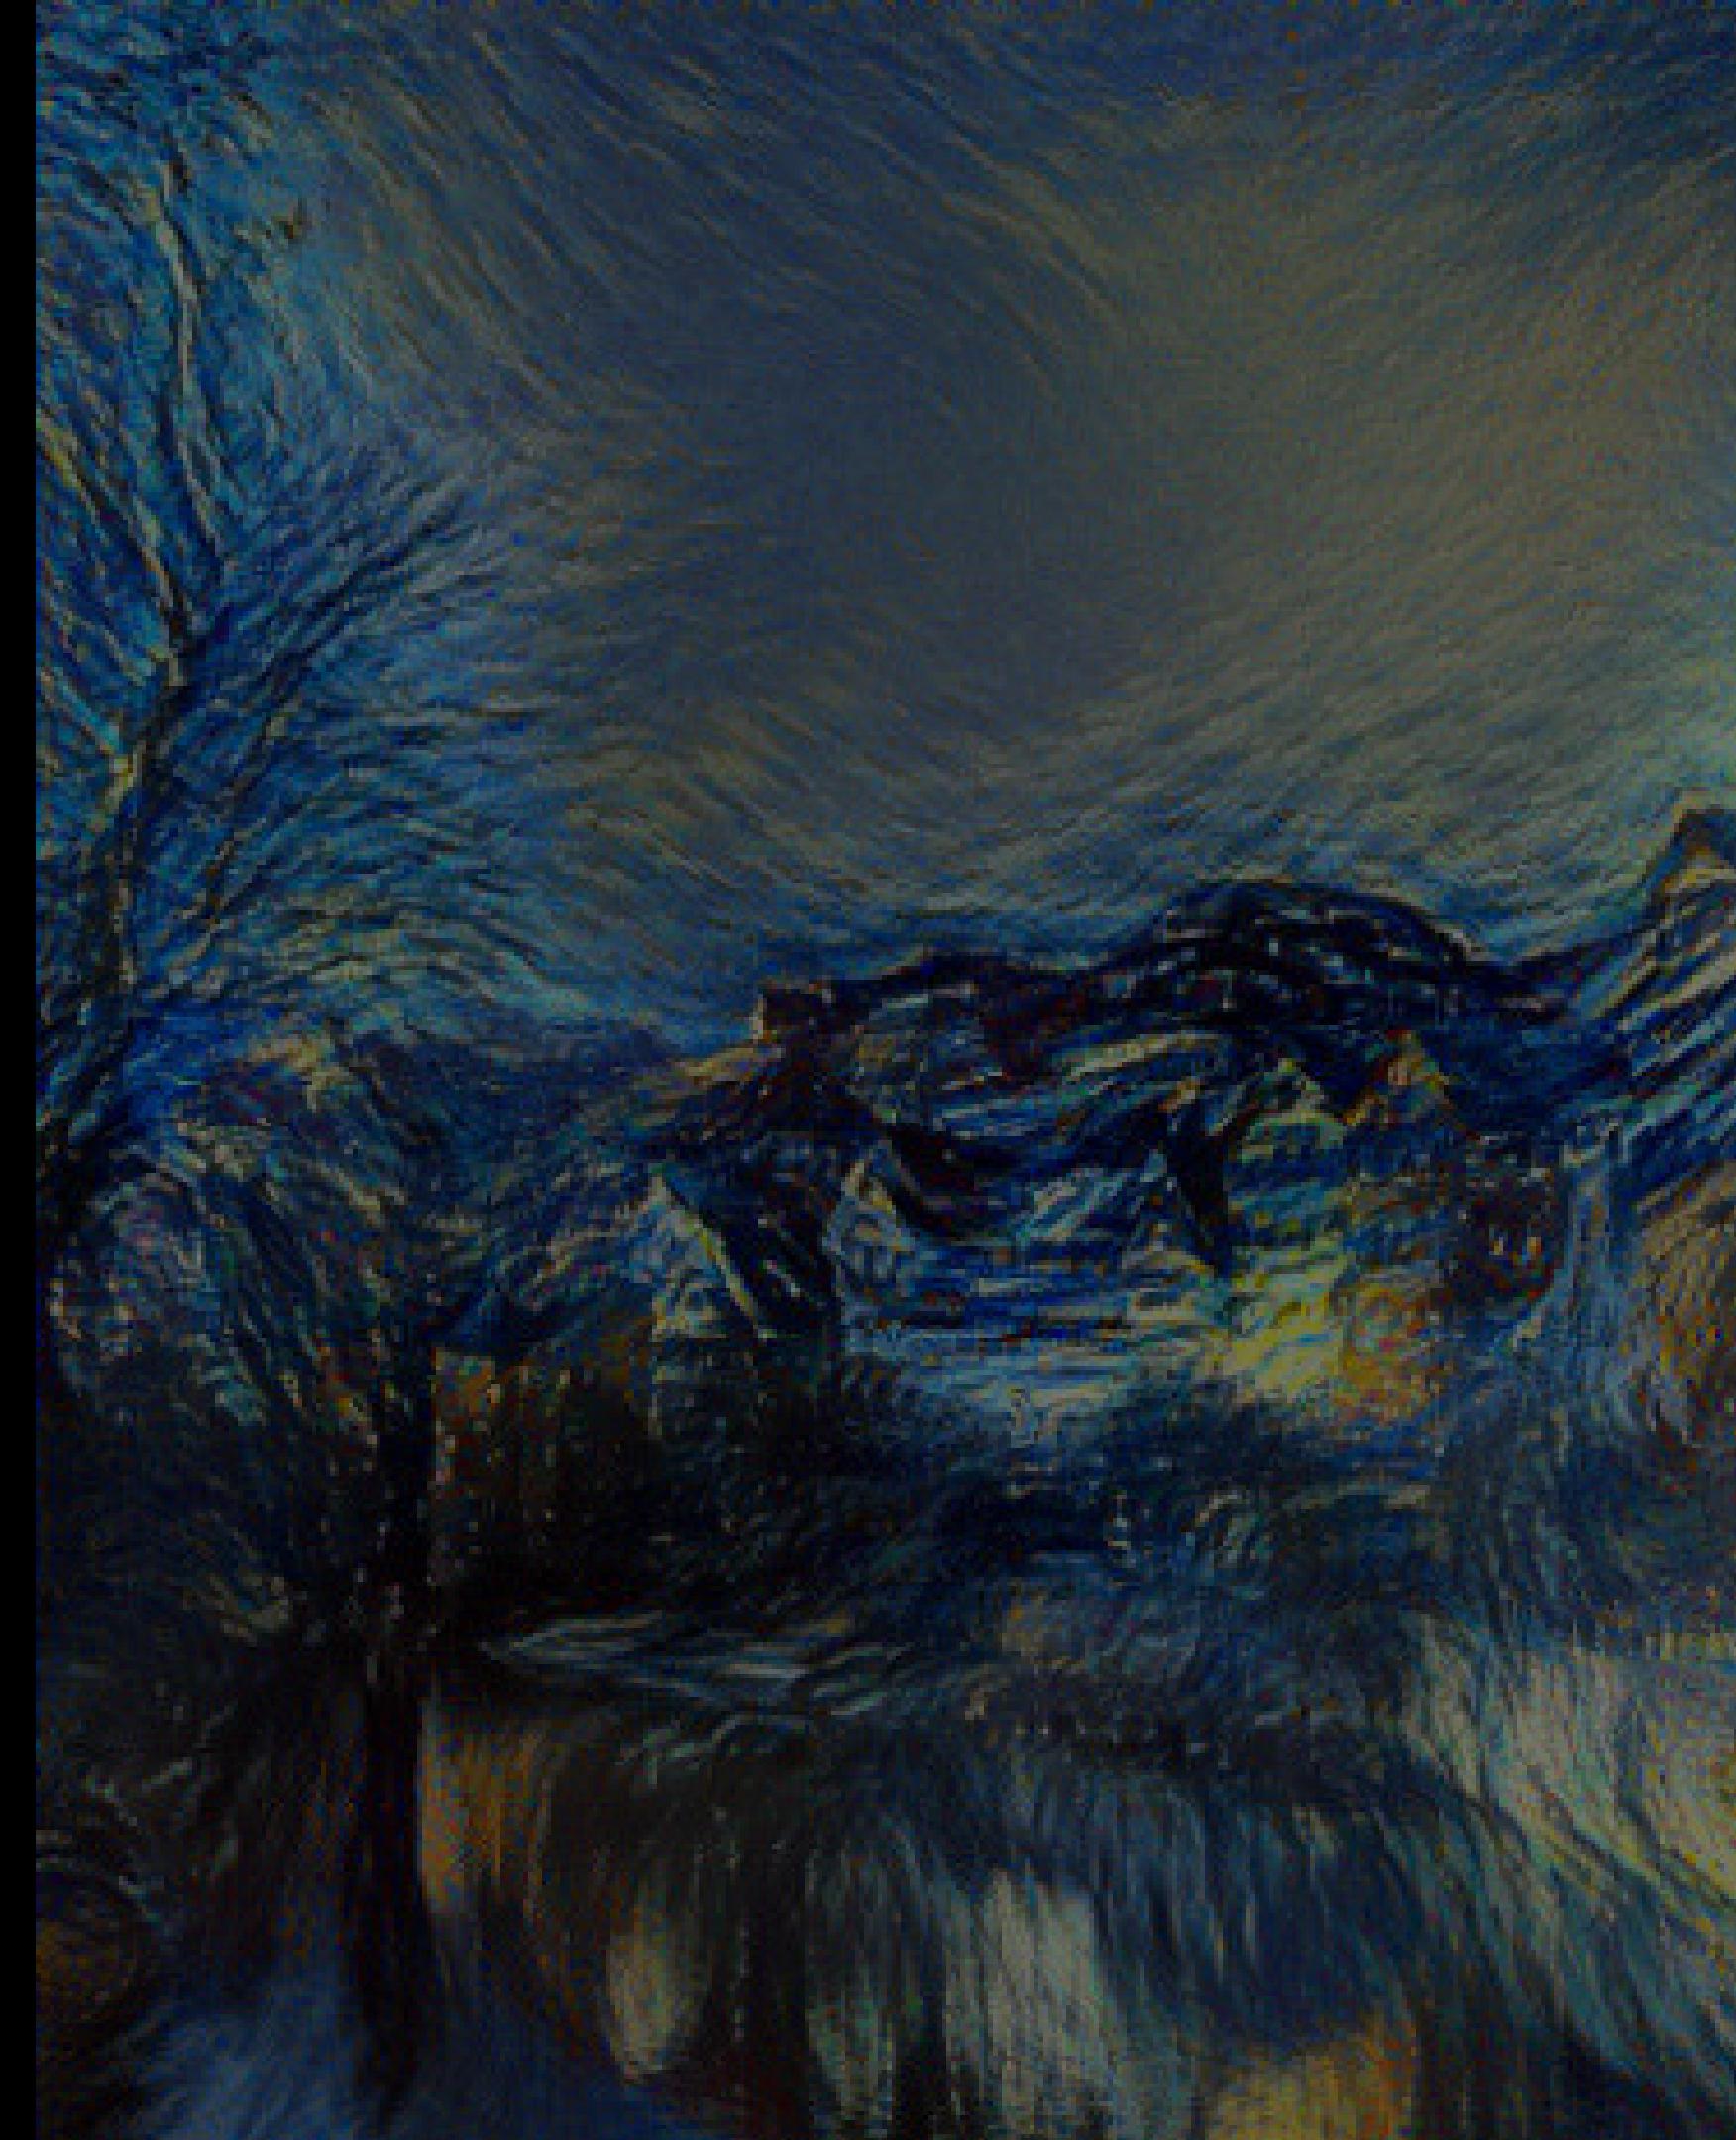


# *The Versatility of the Algorithm*



# *A Recap*

- 1. hierarchies of a pre-trained CNN like VGG19.  
The process is an optimization problem, driven  
by a composite**
- 2. loss function that balances structural  
preservation, stylistic mimicry, and image clarity.  
The choice of optimizer is critical. While Adam  
provides a solid**
- 3. baseline, L-BFGS-B offers a faster path to  
sharper results, especially when implemented  
efficiently with a pattern like the Evaluator class.**





*Thank you  
for your attention*

